

S-Bus: Automatic Read-Set Reconstruction for Multi-Agent LLM State Coordination

Sajjad Khan

Abstract—Concurrent LLM agents sharing mutable natural-language state produce *Structural Race Conditions* (SRCs): write–write and cross-shard stale-read conflicts that silently corrupt agent output. Existing multi-agent frameworks (LangGraph, CrewAI, AutoGen) provide no write-ownership semantics over shared state; conflicts are detected post-hoc, if at all.

We present S-BUS, an HTTP middleware whose central technical mechanism is a server-side *DeliveryLog*: a per-agent log of HTTP GET operations that automatically reconstructs each agent’s read set at commit time without agent SDK changes under HTTP/1.1. The *DeliveryLog* turns ordinary HTTP traffic into a verifiable read-set, enabling optimistic concurrency control over multi-agent shared state with zero in-agent coordination code. The consistency property the *DeliveryLog* provides—we call it *Observable-Read Isolation* (ORI), a partial causal consistency over the HTTP-observable projection of the read set—prevents structural race conditions when agents collaborate via shared shards. We measure this observable projection at 26.1% of single-step references on our principal workload, with session-scoped *DeliveryLog* accumulation extending observable coverage of *self-reported* references to 99.8% over a session—a structural-coverage figure whose denominator is itself contaminated by self-report over-claim (29%–37%, §VII-I); the deflated upper bound on genuine-causal-read coverage is therefore $\leq 70\%$ on our principal workload (PH-2/PH-3 evaluation, gpt-4o-mini, $n=2,100$ step-logs). S-BUS targets the dedicated-shard topology in which each agent owns a distinct write key and reads from shared reference shards.

The paper makes three contributions. **(C1)** We introduce the *DeliveryLog* as a mechanism for automatic HTTP-traffic-based read-set reconstruction in multi-agent LLM systems, and formalise the consistency property it provides (ORI, partial causal consistency over the observable read projection). Mechanised evidence at three tiers: READSETSOUNDNESS and ORICOMMITSAFETY are machine-checked in TLAPS modulo one retained foundational typing axiom (FUNTYPERECONSTRUCTION, §III-D); exhaustive TLC at $N=3$ explores 20,763,484 distinct states to depth 28 with zero violations, and a reduced configuration at $N=4$ explores 2,811,301 distinct states to depth 24 with zero violations; Dafny discharges 9 inductive soundness lemmas (19 verification obligations) on the abstract algorithm. Implementation refinement to the Rust source is empirical, not mechanised. **(C2)** We demonstrate empirical structural-conflict prevention parity against matched-mechanism OCC backends: across PostgreSQL 17 *SERIALIZABLE*, Redis 7 *WATCH/MULTI*, and S-BUS on shared-shard contention sweeps with 427,308 active HTTP-409 conflicts, we observe zero Type-I corruptions (Rule-of-Three upper bound 7.0×10^{-6}). On a non-code workload (Exp. WORKLOAD-B, data-pipeline architecture planning, $n=80$, 8 domains), server-side instrumentation records 0/638 divergent commits under ORI-ON and 590/639 under ORI-OFF ($\chi^2=1,094.98$, $p < 10^{-240}$). **(C3)** We characterise ORI’s operating envelope: in dedicated-shard workloads (Exp. DEDICATED-SHARD, $n=600$) ORI is semantically neutral; in single-shard collaborative writing it is harmful because preservation propagates concurrent contradictions. Both regimes are delimited by paired experiments with prescribed deployment scope; an adaptive merge-routing extension addressing the harmful regime is in separate development.

Scope and honest limitations. The formal guarantees in C1 cover *structural* (Type-I) conflict prevention over the observable read projection; they do not cover *semantic* (Type-II) coherence between concurrent agent outputs. Distributed safety is empirically validated (Exp. DR-9, $n=30$ trials, Wilson 95% CI [0.886, 1.000]) but not TLAPS-mechanised; a ~ 5 ms concurrent-failover window is documented as a known gap. Semantic-quality results rest on an LLM-judge ($\kappa=0.46$ inter-judge agreement on the validation study) without human inter-annotator validation; this is the principal residual evidential gap. Backbone generalisation is established on three vendors (gpt-4o-mini, Anthropic Haiku 4.5, Google Gemini 2.5 Flash) but only for safety parity, not task quality. Workload generalisation is established structurally on two distributions (SWE-bench-derived code coordination and data-pipeline architecture planning) but not for additional non-code classes (document authoring, agent planning, RAG orchestration).

Index Terms—*DeliveryLog*, Automatic Read-Set Reconstruction, Optimistic Concurrency Control, Multi-Agent LLM State, Observable-Read Isolation, HTTP Middleware, Projection-Based Consistency, Rust, TLA+, Dafny.

1 INTRODUCTION

Concurrent LLM agents sharing mutable state produce *Structural Race Conditions* (SRCs):

Definition 1 (Structural Race Condition). *A history H over a shard s contains a Structural Race Condition iff two agents α_i, α_j both read s at version v , generate deltas δ_i, δ_j , and both commit with expected version v without an intervening re-read after the first commit.*

Equivalently: in Adya’s dependency graph [15], H contains a write–write edge (ww) on s with coincident read events but no ordering constraint from a cross-shard validation step. This is a purely syntactic property of the history, decidable from the HTTP-visible event trace; it does not depend on the semantic content of the deltas.

Remark 1 (On “semantic” race conditions). *Prior drafts of this paper defined SRC as producing a final state “outside every valid goal trajectory.” Reviewers correctly flagged this as circular: “valid goal trajectory” has no agreed operationalisation for NL state. Definition 1 is strictly structural and decidable. The semantic effect of a structural race—whether the final state is coherent—is a separate model-dependent question quantified empirically in Exp. DEDICATED-SHARD and Exp. SHARED-STATE, with explicit acknowledgement that an LLM-as-judge rubric is not a substitute for human semantic evaluation (Limitation 5).*

In practice: two agents read the same shard, generate independently valid deltas, and one silently overwrites the other. Existing frameworks (LangGraph [1], CrewAI [2], AutoGen [3]) provide

TABLE 1

Key quantitative findings at a glance. Each row links to the section where it is established. CI columns are Wilson 95% where binomial; “—” where not applicable.

Finding	Value	Where
R_{obs} structural coverage	26.1% of reads (PH-2 workload)	§7.7
R_{hidden} gap (PH-2 workload)	$p_{\text{hidden}} = 0.739$ (CI 0.736–0.741)	§7.7
DL-cumulative observable coverage	99.8% of self-reported refs (note: denom over-claims by 29–37%)	§7.11
R_{hidden} gap (PH-3 workload)	$p_{\text{hidden}} = 0.074$ (CI 0.061–0.081)	§7.8
Semantic extraction recall (PH-3)	0.59 (gpt-4o analyst)	§7.8
Semantic extraction precision (PH-3)	0.92 (gpt-4o analyst)	§7.8
Self-report over-claim rate	29%–37% (2 judges)	§7.9
Inter-judge κ	0.46 (moderate)	§7.9
Type-I corruptions (under active contention)	0 / 427,308 commits	§
Type-I corruptions (full sweep, all conditions)	0 / 884,110 commits	§7.13, §
Workload-B view-divergence (ORI-OFF)	590 / 639 commits divergent	§7.6
Workload-B view-divergence (ORI-ON)	0 / 638 commits divergent	§7.6
TLAPS obligations proved	687 / 687 (modulo 1 retained typing axiom)	§V
TLC state-space ($N=3$, exhaustive)	20,763,484 states (depth 28)	§V
TLC state-space ($N=4$, reduced)	2,811,301 states (depth 24)	§V
TLC Raft (3-node abstract)	247,000 states	§V

no write-ownership semantics for natural-language state. S-BUS addresses this gap with a lightweight HTTP middleware applying OCC to NL agent state, with formal guarantees over the observable read fraction.

1.1 Motivating Example

Four agents collaborate on Django bug #11019 (queryset ordering), each owning a dedicated shard: α_1 owns `orm_compiler`, α_2 owns `migration_script`, α_3 owns `test_fixtures`, α_4 owns `review_notes`. All agents share one read-only reference shard: `db_schema`.

The SRC without S-BUS. Both α_2 and α_3 issue GET /shard/db_schema at $v = 3$ (PostgreSQL dialect). α_1 then commits `db_schema` to $v = 4$ (switching to SQLite for testing). α_2 and α_3 each commit their own shards based on the stale $v = 3$ reading—their output is internally consistent but describes the wrong database. No error is raised; the stale schema silently poisons two shards.

With S-BUS (ORI). When α_2 commits `migration_script`, the ACP validates α_2 ’s `DeliveryLog`: it recorded (`db_schema`, `v3`) at GET time. Since `db_schema` is now at $v = 4$, the ACP rejects with HTTP 409 (CROSSSHARDSTALE). α_2 re-reads, regenerates the migration, and commits successfully. α_3 receives the same rejection and likewise corrects its tests. All four shards converge to a consistent SQLite-based solution.

Key topology. Each agent owns a distinct shard; they read the shared `db_schema` shard but do not write to it concurrently. This dedicated-shard topology is ORI’s primary use case. Single-shard collaborative writing (all agents writing to the same key) requires sequential coordination; Exp. SJ-v4 (§VII-G) and Exp. SHARED-STATE (§VII-M) quantify why.

1.2 Why SWE-bench for Multi-Agent Evaluation?

SWE-bench tasks [40] were designed for single-agent evaluation. We use them for multi-agent coordination evaluation for two reasons: (1) each task has ground-truth acceptance tests, providing an objective semantic correctness criterion independent of the coordination mechanism; (2) tasks are decomposable into role-specialised subtasks that naturally induce shared-state coordination.

We verified decomposability: in 100% (20/20) pilot trials, $N = 4$ agents with distinct shard assignments produced consistent non-contradictory state without structural conflicts. Transfer to other domains (customer service, scientific synthesis) requires independent evaluation; p_{hidden} may differ.

1.3 Scope

S-BUS provides structural conflict prevention for the *dedicated-shard topology*: each agent owns a distinct write key and reads from shared reference shards. Conflict detection operates over the HTTP-observable read fraction R_{obs} ; cross-shard staleness violations are detected and rejected at commit time. Single-shard collaborative writing is out of scope; §11 discusses an adaptive merge-routing extension currently under separate development.

The formal proofs (TLAPS, TLC, Dafny) cover the abstract algorithm’s internal consistency: read-set monotonicity (READSETSOUNDNESS) and cross-shard equality at commit (ORICOMMITSAFETY). They are not proofs about agent semantic correctness, which is workload- and backbone-conditional. The HTTP/1.1 reliance for FIFO-per-connection ordering is stated as Assumption A1; HTTP/2 multiplexing breaks A1 and requires either reverse-proxy pinning or explicit ARSI-mode read-set declaration.

A precise enumeration of what is claimed and not claimed—including the residual coverage gap on the multi-agent workload, the workload distinction between Exp. PH-2 and Exp. PH-3, and the distributed-correctness gap—is integrated into the contribution list (§I-D) and the Limitations section (§8).

1.4 Contributions

This paper makes three technical contributions.

(C1) Automatic HTTP-traffic-based read-set reconstruction with formal guarantees. We introduce the *DeliveryLog*, a server-side per-agent log of HTTP GET operations that automatically reconstructs each agent’s read-set at commit time. The *DeliveryLog* turns ordinary HTTP traffic into a verifiable read-set, allowing optimistic concurrency control to be applied to multi-agent shared state without requiring agents to declare what they read or write

coordination code. Architecturally, the DeliveryLog is a scope-restricted COPS-style causal log [16] adapted from distributed-database literature to the multi-agent LLM setting. We formalise the consistency property the DeliveryLog provides and call it *Observable-Read Isolation* (ORI, Definition III.4): a partial causal consistency over the HTTP-observable projection of an agent’s read set. We position ORI in Adya’s isolation lattice [15] and establish via counterexample histories that, when projected to the observable read set, ORI sits strictly between Read-Committed and Snapshot Isolation: ORI prevents G2-item which RC does not, and permits write-skew on the unobservable fraction which SI does not. Three tiers of mechanised evidence support the model: TLAPS module `SBus_TLAPS_v16.tla` discharges 687 obligations (zero failed, modulo one retained foundational typing axiom FUNTYPINGRECONSTRUCTION; §III-D) proving READ-SET-SOUNDNESS (recorded-read monotonicity) and ORICOMMIT-SAFETY (cross-shard equality at commit time, directly capturing Definition III.4(2)) for arbitrary N_{agents} ; exhaustive TLC at $N=3$ explores 20,763,484 distinct states to depth 28 with zero invariant violations, and a reduced configuration at $N=4$ explores 2,811,301 distinct states to depth 24 with zero invariant violations; Dafny module `sbus_lemmas_v4.dfy` machine-checks 9 inductive soundness lemmas (19 verification obligations). Implementation refinement to the Rust source is empirical, not mechanised.

(C2) Empirical safety parity with production OCC. Exp. PG-COMPARISON implements the multi-agent workload against three independent OCC engines: S-BUS (Rust), PostgreSQL 17 `SERIALIZABLE` (`pg_sbus_server.py`), and Redis 7 `WATCH/MULTI` (`redis_sbus_server.py`). Across $N \in \{4, 8, 16, 32, 64\}$, 30 task domains, 1,350 runs, and 200,880 commit attempts, all three backends exhibit zero Type-I corruptions (95% Rule-of-Three upper bound 1.49×10^{-5}). The contention extension (Exp. PG-CONTENTION) adds 472,750 commit attempts under shared-shard contention with 427,308 active HTTP-409 conflicts and zero Type-I corruptions; SCR agreement across the three backends is within 1 pp at $N \geq 8$. Cross-backbone paired replication on Anthropic Haiku 4.5 and Google Gemini 2.5 Flash ($n=2,400$ each) confirms safety parity across vendors. The architectural value of S-BUS over transactional-DB baselines is operational simplicity and the LLM-native contract, not throughput.

(C3) Topology-conditional operating envelope. The ACP’s core invariant—preservation of every commit’s contribution—is unconditional structurally and topology-conditional semantically. In 959 paired trials of Exp. ORI-ISOLATION, ORI-ON preserves 40/40 agent-step contributions per trial; ORI-OFF (last-writer-wins) preserves 10/40 = $1/N$. Whether this preservation is beneficial depends on workload topology. In dedicated-shard workloads (Exp. DEDICATED-SHARD, $n=600$): ORI is semantically neutral—100% coherent in both fresh and stale conditions. In single-shard collaborative-writing workloads (Exp. SHARED-STATE, $n=180$): ORI is semantically harmful—100% contradicted under ORI-ON vs. 85.6% under ORI-OFF, because preserving all contributions preserves their mutual contradictions. We prescribe deployment scope accordingly (Box 2). Single-shard collaborative writing is the natural extension target: the structural preservation guaranteed by ORI is not the right primitive for that regime, and an adaptive merge-routing protocol is required (§11).

Workload-scope disclaimers: Two empirical claims carry workload-scope qualifications which we state once here and refer back to throughout. First, the observable read fraction R_{obs} covers 26.1% of reads on the multi-agent workload (Exp. PH-2); the

remaining $p_{\text{hidden}} = 0.739$ is not directly observable at the HTTP layer within a single step. Section III-D decomposes this residual structurally and identifies session-scoped DeliveryLog accumulation as the dominant coverage mechanism. Second, semantic-extraction evaluation (Exp. PH-3, 0.59 recall / 0.92 precision) was conducted on a single-agent rotating-target workload at $p_{\text{hidden}} = 0.074$; transfer to the high- p_{hidden} multi-agent regime is open (§11). The TLAPS theorems are proofs about the state machine’s internal consistency over R_{obs} , not proofs about agent correctness over the full read set.

Distributed scope: P1 session replication is validated empirically in Exp. DR-9 (30/30 ORI invariants survived leader failover). A residual ~ 5 ms concurrent-failover window within the fire-and-forget replication interval is the remaining gap; full Raft-TLAPS mechanisation is future work (Limitation 11).

2 RELATED WORK

We position S-BUS against four bodies of prior work: multi-agent LLM frameworks (which provide the deployment context but lack write-ownership semantics), formal verification of distributed systems (the methodology S-BUS’s three-tier evidence draws on), classical concurrency control (the technical foundation S-BUS adapts), and isolation-level theory (which gives us a precise position to claim). We close with a “Why not X?” subsection addressing the most natural alternatives.

2.1 Multi-Agent LLM Frameworks

Production multi-agent frameworks route shared state through workflow graphs (LangGraph [1]), message-passing channels (AutoGen [3]), or agent-local memory (CrewAI [2]). None provide write-ownership semantics over mutable shared state; conflicts are detected post-hoc via agent self-reports or downstream validation, if at all. Other frameworks (MetaGPT [4], CAMEL [5], Swarm [6], Agent-to-Agent [7], ReAct [8], DSPy [9], Semantic Kernel [10], AgentScope [51], Voyager [52], SWE-agent [53]) make analogous design choices.

The empirical consequence is documented by Cemri et al. [11], who taxonomise multi-agent LLM failure modes across 200 production traces and find that consistency errors account for 23–31% of LangGraph failures. We replicate this on a controlled microbenchmark (§7.20): with $N=4$ concurrent agents writing to a shared key, 20/20 trials produce silent overwrites under last-write-wins. S-BUS targets exactly this category, providing the structural primitive (ORI) the existing frameworks lack.

The closest related work is Park et al.’s “Generative Agents” [12], which addresses NL coordination via a shared memory stream with reflection-based abstraction. Park et al. explicitly note their architecture would benefit from “stronger consistency guarantees on the memory stream” but do not propose a mechanism. ORI is complementary: it provides the structural guarantee Park’s memory stream lacks. Production memory systems (Letta [48], MemGPT [14], Zep, Mem0, LangMem) assume single-writer-per-item and provide no cross-agent transactional semantics; they are orthogonal to S-BUS—these systems provide storage and context-window management, S-BUS provides the consistency layer.

Session guarantees: Terry et al. [13] introduced session guarantees (read-your-writes, monotonic reads, writes-follow-reads) for Bayou. ORI’s DeliveryLog enforces an analogous per-session causal ordering for HTTP-observable reads, making ORI a

partial causal consistency model in the spirit of Bayou, restricted to R_{obs} .

2.2 Formal Verification for Distributed Systems

The state of the art for formally verified distributed systems sits on a spectrum from full implementation refinement to algorithm-level TLA+ specifications. IronFleet [25] occupies the rigorous end with full Dafny refinement at 10+ person-years of proof engineering. Verdi [26] provides the first machine-checked proof of Raft requiring 90+ invariants. Closer to industrial practice, TigerBeetle [28] and FoundationDB [27] validate concurrency control via deterministic simulation testing without TLA+-to-implementation refinement; TiDB, CockroachDB, and etcd publish TLA+ specifications without implementation refinement proofs.

S-BUS’s three-tier evidence sits within this spectrum: TLC exhaustive model-checking for $N \leq 4$, TLAPS for the abstract algorithm at arbitrary N , and Dafny inductive lemmas on types structurally equivalent to the Rust implementation. Full Rust refinement via Verus [54] or Creusot [55] is future work; both toolchains reached usable status in 2023–2024 for synchronous Rust, but Verus’s async support (required for S-BUS’s tokio-based implementation) remains under active development.

2.3 Concurrency Control

S-BUS adapts the OCC tradition (MVCC [33], OCC [34], STM [35], TL2 [36], Calvin [37], Percolator [38]) for the multi-agent LLM domain. We focus this discussion on the four pieces of prior work whose design choices map most directly onto S-BUS’s: FoundationDB’s Directory Layer for the per-shard observable-state model; Cherry-Garcia for the bolt-on architecture; HTTP RFC 7232 for the conditional-request mechanism; and COPS for the causal-log methodology.

FoundationDB Directory Layer: The FoundationDB Directory Layer [47] provides cross-key serializable transactions over opaque byte strings via hybrid logical clocks, the closest structural analogue to S-BUS’s per-shard model in a production database. S-BUS differs in that the read-set is reconstructed automatically from HTTP GET traffic via the DeliveryLog rather than declared at transaction-begin time; this is the property that makes the LLM-agent use case viable without per-agent transaction scoping. CockroachDB [39] and Spanner [20] use similar HLC and TrueTime mechanisms respectively for distributed serializable OCC, demonstrating that S-BUS’s per-key OCC model is distribution-friendly.

Cherry-Garcia: bolt-on transactions: Cherry-Garcia [50] layers serializable transactions on heterogeneous NoSQL key-value stores without server-side modification, using a client-side protocol to coordinate cross-store writes. S-BUS’s architecture is structurally similar: transactional semantics layered over a non-transactional substrate (HTTP GET/POST against an in-memory registry), with the DeliveryLog playing the role of Cherry-Garcia’s client-side read-set tracker. The principal difference is observation scope: Cherry-Garcia’s client declares the transaction’s read-set explicitly; S-BUS’s server reconstructs it from HTTP-observable GET traffic, enabling SDK-free integration under HTTP/1.1.

HTTP conditional requests (RFC 7232): The IETF HTTP/1.1 conditional-request mechanism [44] (If-Match with ETag values) has provided per-resource optimistic concurrency control since 1999. S-BUS’s ACP is an adaptation of this pattern with two specific extensions: (i) automatic read-set reconstruction

via the DeliveryLog (RFC 7232 leaves read-set tracking to the client), and (ii) cross-shard validation—an agent’s commit to key k is aborted if the recorded version of a sibling key k' the agent previously read has since advanced. Reviewers who read the system as “ETags + DeliveryLog” are substantively correct; the cross-shard reconstruction mechanism and the empirical evidence that it matters under LLM agent workloads are the contributions.

COPS and Eiger: causal consistency: COPS [16] orders writes by causal dependency. The DeliveryLog captures the causal chain from GET to commit within a session, with cross-shard validation enforcing that causally preceding reads have not been superseded. Differences from COPS are scope and representation: COPS tracks all writes system-wide via vector clocks; DeliveryLog tracks per-session HTTP GETs only (26.1% empirically), using scalar (k, v) pairs sufficient for OCC validation. The DeliveryLog is therefore not a new primitive but a restricted causal log sufficient for OCC validation of the observable fraction.

Single-node OCC pioneers: Silo [29], Hekaton [30], FaRM [31] establish the performance envelope for single-machine OCC on structured data; Boki [56] provides a serverless analogue. S-BUS’s per-key mutex implementation is architecturally simpler because NL deltas are not amenable to the epoch-based batching those systems rely on; throughput is bounded by independent shard count, not CPU-local epoch advance. Microsoft Orleans [32] confines mutable state to single actors, providing implicit ownership; S-BUS’s Ownership Token formalises the same invariant per-shard, with cross-shard stale-read prevention added via the DeliveryLog (no equivalent in actor frameworks).

2.4 Isolation Levels and Anomaly Taxonomy

Following Adya [15], ORI occupies a specific position in the dependency-graph hierarchy: it prevents G0, G1a, G1b, and G2-item over R_{obs} , permits G2 (write-skew) over R_{hidden} , and permits G3 (phantom) globally. This is a projection-based consistency model: correctness guaranteed over the observable projection of the history (Table 2). RedBlue consistency [17] provides a theoretical parallel: HTTP GETs are analogous to blue operations (no coordination), ACP commits are red (serialised); ORI can be read as “the red operation validates all blue operations in its read-set at commit time.” The gap between ORI and full SI is that R_{hidden} reads are neither red nor blue—they are invisible to the coordination layer entirely. Parallel Snapshot Isolation [18] is a weakening of SI for geo-distribution; ORI is further weakened in that it only validates R_{obs} reads. SSI [19] extends SI to full serializability via anti-dependency detection; ORI sits strictly below SSI and SI on any read-set comparison over R_{obs} . RAMP [21] provides multi-key read-atomic consistency without a single-node coordinator; S-BUS’s cross-shard stale-read rejection provides an analogous guarantee, but RAMP requires typed schemas and pre-declared read-sets which S-BUS replaces with DeliveryLog reconstruction.

Coordination avoidance: Bailis et al. [22] show coordination can be avoided when operations are I-confluent. NL writes are non-I-confluent: “use PostgreSQL” and “use DynamoDB” are not commutative under a type-consistency invariant, formally justifying OCC as the required coordination mechanism for NL state.

2.5 LLM-Assisted Merge vs. OCC

LLM-assisted merge is a legitimate alternative: rather than aborting conflicting NL deltas, an LLM resolves conflicts post-hoc. Exp. MERGE (Arc A5) compares both approaches empirically.

For additive shards, character-level CRDTs (Automerger [43], Yjs, Loro [42]) handle concurrent edits without OCC overhead and are preferable. S-BUS’s OCC targets the non-commutative case where no CRDT merge function is well-defined. Loro’s movable-tree CRDT and Automerger’s rich-text CRDT handle some non-commutative operations via operation-intent preservation rather than last-writer-wins; these are compatible with S-BUS for additive shards but do not obviate OCC for mutually-exclusive NL state.

2.6 Why Not X?

This subsection addresses the most natural alternatives a reader may consider before reading further.

Why not Raft for the registry?: Raft would serialise all commits through a single log, providing linearizability across all keys. S-BUS’s per-key OCC achieves linearizability per-key without the cross-key serialisation cost: two agents committing to disjoint keys are not serialised. For the multi-agent LLM workload, where agents typically own disjoint shards, this distinction matters. S-BUS’s distributed extension uses Raft for log replication of metadata (P1 session replication, Exp. DR-9) but retains per-key OCC for commits.

Why not CRDTs?: CRDTs require a commutative merge function. NL state is not universally commutative: “schema uses PostgreSQL” and “schema uses DynamoDB” have no merge function preserving both. CRDTs work for the subset of NL state that is additive (append-only logs, character-level edits), and S-BUS is compatible with such shards. For the non-additive case—which dominates structured technical state—OCC is the appropriate primitive. The adaptive-routing extension (§11.1) addresses workloads where both regimes coexist on the same shard.

Why not application-level locks?: A naive solution: each agent acquires a process-level mutex before issuing a commit. This works but serialises all commits, $N \times$ -amplifying latency for N concurrent agents. S-BUS’s per-key OCC permits parallel commits to disjoint keys with rejection on cross-shard staleness, achieving the same correctness with $O(1)$ latency per commit. Empirically (Exp. SEQUENTIAL, Arc A1), the wall-time difference is 5–10 $\times$ at $N=5$.

Why not Temporal/durable execution?: Temporal.io [46] and Golem provide exactly-once workflow execution with persistent state and automatic retry. Temporal’s recent Update API adds synchronous workflow-update primitives architecturally similar to ORI’s commit validation, but per-workflow rather than per-shard; the cross-shard read-set inference ORI gets from the DeliveryLog is not a feature of Temporal Updates. S-BUS targets the case where N agents run concurrently on shared NL state without a workflow orchestrator.

Why not ETags directly?: RFC 7232’s ETag mechanism provides per-resource OCC. For S-BUS’s target use case (cross-shard stale reads), ETag is insufficient because the agent’s read of shard k' is not part of the commit to shard k . The DeliveryLog reconstructs the cross-shard dependency server-side; ETags do not. We are explicit (§II-C) that S-BUS is best read as “ETags + DeliveryLog”; the contribution is the DeliveryLog’s automatic read-set reconstruction.

Why not snapshot isolation with read timestamps?: Spanner [20] and CockroachDB [39] provide snapshot reads at a chosen timestamp via TrueTime or HLC, which would let agents read a consistent cross-shard snapshot and commit against that

TABLE 2
Anomaly classes: ORI prevents vs. permits (Adya [15]).

Anomaly	Status	Reason
G0 (dirty write)	Prevented	Global write lock
G1a (dirty read)	Prevented	Version check
G1b (intermediate)	Prevented	Atomic commit
G2-item (anti-dep.)	Prev. over R_{obs}	DeliveryLog
G2 (write skew)	Perm. for R_{hidden}	Out of scope
G3 (phantom)	Permitted	No predicate tracking
Type-III (semantic)	Permitted	Structural $\neq$ goal

timestamp. This is a viable alternative architecture but requires either (a) time-synchronised hardware, which most multi-agent deployments do not have, or (b) hybrid logical clocks with explicit per-write coordination, which adds latency to every read. S-BUS trades the snapshot guarantee for a weaker per-shard guarantee with substantially simpler deployment: agents use ordinary HTTP, and the DeliveryLog reconstructs the read-set from observed GET traffic without coordinating timestamps. For deployments where snapshot reads are available (e.g., a Spanner-backed registry), agents can layer ARSI-mode reads on top of a snapshot endpoint to recover SI semantics; this composition is straightforward but beyond the scope of this paper.

Why not session-typed channels or coordination middleware?: Session types [57] encode coordination invariants statically in agent code, and ZooKeeper / etcd provide strongly-consistent coordination primitives via watch-based notification. Both target the case where coordination requirements are knowable at design time. S-BUS targets the opposite case: agents are heterogeneous black-box LLM workers whose access patterns emerge at runtime from natural language prompts. Static session typing is therefore infeasible (you cannot type-check an LLM’s emitted GETs); ZooKeeper-style watches are too coarse-grained (the unit of consistency is a key, but agents read multiple keys per step). A hybrid deployment using S-BUS for per-shard OCC and ZooKeeper for coarser-grained agent-coordination metadata is plausible and not in conflict with this paper’s claims.

3 MECHANISM AND CONSISTENCY MODEL

This section presents S-BUS’s central technical contribution—the *DeliveryLog* mechanism for automatic HTTP-traffic-based read-set reconstruction—and then formalises the consistency property it provides, which we name *Observable-Read Isolation* (ORI). We present mechanism and property in this order because the DeliveryLog is the engineering novelty: it is what allows the system to apply optimistic concurrency control over an agent’s read-set without requiring the agent to explicitly declare what it read. ORI is the formal characterisation of what guarantees this mechanism provides over the resulting observable projection.

The DeliveryLog at a glance: For each agent α , the server maintains a $\text{DeliveryLog}_\alpha$: an append-only log of (k, v) pairs recorded whenever α issues an HTTP GET for shard k and the server returns version v . At commit time, the DeliveryLog serves as the agent’s reconstructed read-set: the server compares each (k', v') in $\text{DeliveryLog}_\alpha$ to the registry’s current version of k' , and rejects the commit if any cross-shard version has advanced. This is HTTP ETags generalised across shards, with the read-set reconstructed server-side from observed traffic rather than declared by the agent. The novelty is not the mechanism in isolation (R_{obs} -style read-set

tracking is well-known); the novelty is that the agent does not declare anything and writes no coordination code. Sections 3.4 and V formalise the consistency property ORI that this mechanism provides.

3.1 Definitions

Definition 2 (Shard). A key-value pair (k, v, c) : key k , version $v \in \mathbb{N}$, content $c \in \Sigma^*$ (opaque NL string).

Definition 3 (Agent Read Set). $R_\alpha = R_{obs} \cup R_{hidden}$: observable reads are HTTP GET calls; hidden reads are shard-key references in conversation history.

Definition 4 (Effective Read Set). $\hat{R} = R_{explicit} \cup \text{DeliveryLog}_\alpha$. Under Assumption III.1, $\hat{R} \supseteq R_{obs}$.

Remark 2 (Ownership token and collaborative writes). The ownership token model enforces per-shard exclusive writes: only the agent currently holding the token can commit to a given shard. This prevents write-write conflicts but also prevents two agents from simultaneously contributing to the same shard. Collaborative writes to the same shard require explicit agent-level serialisation (sequential handoff) or ARSI mode with declared read-sets. Workloads where agents naturally own disjoint shards (the Django #11019 case study, §VII-O) are the intended use case.

Definition 5 (ORI-Legal History). A history H over shards S and agents A is ORI-legal iff:

- 1) **Write-write serialisation.** For every shard $s \in S$, all committed writes to s are totally ordered; no two concurrent commits succeed on the same shard without one being serialised before the other.
- 2) **Cross-shard R_{obs} freshness.** For every committed operation $\text{Commit}(\alpha, k, v)$ and every entry $(k', v') \in R_{obs}(\alpha)$ recorded in α 's DeliveryLog with $k' \neq k$: no committed write to k' appears in H strictly between the recording of (k', v') and $\text{Commit}(\alpha, k, v)$ in the serialisation order.

Forbidden: G2-item [15] projected to R_{obs} . Permitted outside scope: G2-item projected to R_{hidden} , G3 (phantom), Type-III (semantic divergence).

Remark 3 (ORI as π -causal consistency). Define π -causal consistency for projection π as: for every pair of operations $o_1 \xrightarrow{hb} o_2$ (Lamport happens-before [24]) with $o_1, o_2 \in \pi(H)$, every process observing o_2 has observed o_1 . ORI instantiates π -causal consistency with $\pi = R_{obs}$. Standard causal consistency is the special case $\pi = \text{identity}$. ORI is therefore strictly weaker than full causal consistency and is incomparable to points in Adya's isolation lattice defined over the full read set. All lattice-positioning language in this paper is qualified by the projection, e.g. " $RC \prec_{R_{obs}} SI$."

Remark 4 (Composition scope). ORI is defined for a single S-BUS instance over a single agent session. Composition of two ORI-compliant subsystems is not guaranteed to yield ORI: cross-subsystem reads are captured in neither DeliveryLog, which is the R_{hidden} problem applied at the subsystem boundary. Multi-subsystem deployments must use ARSI mode with explicit cross-subsystem read-set declaration.

Remark 5 (Conflict-proximate hypothesis: status). We conjecture that R_{obs} reads are conflict-proximate (issued shortly before dependent commits). Partial evidence: Exp. ORI-ISOLATION

shows $GET \rightarrow COMMIT$ co-location within a single execution step in 100% of 959 trials. We do not treat this as validation of the hypothesis on unstructured or long-horizon workloads; it is a harness artifact of the SWE-bench-style task structure used in these experiments. A GET-to-commit-gap measurement on diverse workloads is pending.

3.2 Assumptions

Assumption 1 (DeliveryLog Completeness (A1)). Every HTTP GET /shard/:key?agent_id=X is recorded in the DeliveryLog exactly once before the response returns. Holds unconditionally on a single node; the Rust implementation holds the same Mutex guard across both the DeliveryLog write and the HTTP response. Scope conditions: (a) buffering: proxies must use proxy_buffering off; (b) streaming: S-BUS records the initial GET, capturing the version at read time; (c) retries: the DeliveryLog entry is recorded at request receipt, not at client acknowledgement; (d) HTTP/2: request reordering under HTTP/2 multiplexing can break the ordering assumption; agents on HTTP/2 should use ARSI mode (Table II).

Assumption 2 (TTL $\geq T_{\max}$ (A2)). Session TTL is set above the experiment wall time.

Assumption 3 (Serialised commit log (A3)). All commits are totally ordered via a single serialisation point: a global write lock on a single-node deployment, or the Raft leader's log under the 3-node distributed deployment (§IX). Raft leader election uses randomised timeouts (configured [500, 1000] ms), making A3 a probabilistic rather than strictly deterministic guarantee. Split-brain is bounded by Raft's safety proof [23] but is not zero under arbitrary network partitions.

Assumption 4 (Observation scope (A4)). $f_{obs} = |R_{obs}|/|R_\alpha| = 1 - P_{hidden}$.

Assumption 5 (Cooperative agents (A5)). Agents do not forge commit requests with fabricated read-sets or version numbers. Byzantine/malicious agents are out of scope. In open multi-agent deployments (e.g., user-provided tools in a marketplace), this assumption may be violated. The production mitigation is HMAC-based request signing: each agent receives a per-session shared secret, and the ACP verifies the HMAC over (k, v_e, δ, α) before processing commits. HMAC authentication is a planned extension; the current implementation targets single-tenant trusted-agent deployments.

3.3 ORI Safety Property

Property 1 (ORI Safety). Under A1–A5: all committed histories are ORI-legal over R_{obs} .

Proof status:

- TLC-verified for $N \leq 4$, $|\text{Shards}|=3$ under the extended spec `SBus_ori.tla` with explicit `READSETSOUNDNESS` and `NOSTALECROSSSHARD` invariants: exhaustive at $N=3$ explores 20,763,484 distinct states to depth 28, and a reduced configuration at $N=4$ explores 2,811,301 distinct states to depth 24. Zero invariant violations at any configuration tested.
- TLAPS-proved for arbitrary N_{agents} : the `v1` artifact (`SBus_TLAPS.tla`) mechanises type safety, ownership uniqueness, and version monotonicity (103 obligations). The `v16` artifact (`SBus_TLAPS_v16.tla`) mechanises

two theorems: `READSETSOUNDNESS` (recorded-read monotonicity) and `ORICOMMITSAFETY` (cross-shard equality at commit, directly capturing Definition III.4(2) as $\forall(k', v') \in dlog[\alpha]$:

$k' \neq k \Rightarrow registry[k'].v = v'$). 687 obligations proved by `tlapm 1.5`, 0 failed. Two sequence-theoretic library facts are discharged mechanically via `SequenceTheorems.SeqDef` and `SequenceTheorems.ElementOfSeq`. One mathematical AXIOM is retained (`FUNTYPINGRECONSTRUCTION`: a function with domain S and values in T is in $[S \rightarrow T]$) as a primitive fact about TLA+'s typed-function-space construction; this is not present in the standard `FunctionTheorems.tla` library (which covers bijections, injections, surjections, and Cantor-Bernstein). Two parameter ASSUMES (`NOOWNERAGENTS`; initial shard content is a `STRING`) are standard TLA+ parameterisation, not mathematical axioms. Together this closes the reviewer critiques “the invariant is weaker than the claim” (the monotonicity property `READSETSOUNDNESS` is a necessary condition; `ORICOMMITSAFETY` captures the full property) and “the four axioms are proof-engineering shortcuts” (three of the v10 axioms are now library-discharged; the one retained is a documented primitive of TLA+'s function-space theory).

- Dafny 4 machine-checks 9 inductive soundness lemmas in `sbus_lemmas_v4.dfy` (19 verification obligations discharged, 0 errors): `INITSOUNDNESS`, `READPRESERVESOUNDNESS`, `COMMITPRESERVESOUNDNESS`, `TIMEOUTPRESERVESOUNDNESS`, `MONOTONICCOMMITPRESERVESOUNDNESS`, `CROSSSHARDSTALENESSISSTRICT`, `OWNERSHIPINVARIANTINDUCTIVE`, `VERSIONMONOTONICITYLEMMA`, `ACPLockOrderIsDeadlockFree`. The earlier vacuous `CrossShardSafetyLemma` in `sbus_lemmas.dfy` is deprecated.

Full-path mechanised proof of the distributed (Raft) case remains future work (Limitation 18); the TLC-checked abstraction is documented in §3.4.

Lemma 1 (DeliveryLog R_{obs} Happens-Before Soundness). *Under A1: for every agent α and every pair of operations $o_1 \xrightarrow{hb} o_2$ where o_1 is an HTTP GET and o_2 is a COMMIT by α , $DeliveryLog_\alpha$ contains the entry (k, v) for the shard read by o_1 before o_2 executes.*

Sketch. A1 guarantees the Mutex guard is held across the Delivery-Log write and HTTP response for every GET. Therefore the entry $(k, reg[k].v)$ is recorded before the response returns to the agent. Since o_2 (COMMIT) arrives after o_1 's response (TCP ordering), $DeliveryLog_\alpha$ contains the entry at commit time. ACP lines 1–4 then validate all entries in $\hat{R} \supseteq R_{obs}$, enforcing the happens-before edge.

3.4 Formal Evidence

We mechanise ORI safety at three tiers; implementation refinement from the Rust source is empirical rather than mechanised, consistent with standard practice short of IronFleet [25].

TLAPS theorems (arbitrary N_{agents}): Module `SBus_TLAPS_v16.tla` states two theorems. `READSETSOUNDNESS` is a state invariant: no agent's recorded read can be ahead of any committed version, $\forall a, i : dlog[a][i].v \leq registry[dlog[a][i].k].v$.

`ORICOMMITSAFETY` is a transition property: the Commit action fires only in states where cross-shard recorded reads match the current registry versions exactly, capturing Definition III.4(2). `tlapm v1.5` closes 687 obligations with 0 failed, covering Init, Read, Commit, Timeout, and inductiveness. Two sequence-theoretic facts are discharged via the standard `SequenceTheorems` library (`SeqDef`, `ElementOfSeq`); two parameter ASSUMES on unspecified constants are standard TLA+ parameterisation, not mathematical axioms.

Retained axiom: an honest accounting: One AXIOM is retained: `FUNTYPINGRECONSTRUCTION`, the statement that if $DOMAIN f = S$ and $\forall x \in S : f[x] \in T$ then $f \in [S \rightarrow T]$. This is the converse of the typed-function-space introduction rule and is foundational to TLA+'s function-space construction. We have not discharged it, and we want to be precise about why. The standard `FunctionTheorems.tla` library covers bijections, injections, surjections, and Cantor-Bernstein, but does not include `FUNTYPINGRECONSTRUCTION` as a derived theorem; attempts to derive it within `tlapm` from the underlying TLA+ function-space axioms have not closed in our hands. The fact is widely treated as obvious in TLA+ practice (it appears as an unchecked side-condition in numerous published proofs), but “widely treated as obvious” is not the same as “mechanically verified.” We therefore treat the proof as machine-checked modulo this single axiom, distinct from Verdi's network-model assumptions (which are about external phenomena outside the formal system) and distinct from IronFleet's zero-axiom discipline (which mechanises every foundational fact).

The practical reader should be aware that if `FUNTYPINGRECONSTRUCTION` were false—which we believe it is not—then the proof's type-correctness layer would not hold, and the higher safety theorems that depend on type-correctness (`READSETSOUNDNESS` and `ORICOMMITSAFETY`) would lose their guarantee. We consider this a low-probability concern given the axiom's foundational nature and its widespread implicit use in published TLA+ proofs, but we do not claim zero-axiom discipline, and we do not want a reviewer to read past this gap. The concrete next step is discharging the axiom via the TLAPS theorem-proving manager's Isabelle/TLA backend, which encodes a deeper layer of TLA+'s set theory than `tlapm`'s default backend reaches; this should either close the obligation or expose precisely what additional foundational fact is needed. The full proof script (`proofs/SBus_TLAPS.tla`, lemma chain L1–L9, axiomatisation rationale) is in the `sbus-formals` repository.

TLC at $N \leq 4$: Module `SBus_ori.tla` adds `READSETSOUNDNESS` as an explicit state invariant alongside `TYPEINVARIANT`, `OWNERSHIPINVARIANT`, and `VERSIONMONOTONICITY`. Zero violations at every configuration tested. The exhaustive $N=3$ run explores 20,763,484 distinct states at depth 28. A reduced configuration at $N=4$ (`MAXVERSION=2`, `RETRYBUDGET=2`) completes at 2,811,301 distinct states at depth 24. A full exhaustive $N=4$ sweep at `MAXVERSION=3` is open work and is not part of the v1 artifact.

Dafny: 9 inductive soundness lemmas: `sbus_lemmas_v4.dfy` machine-checks that `READSETSOUNDNESS` is preserved by every algorithm action (Init, Read, Commit, Timeout, MonotonicCommit) plus `CROSSSHARDSTALENESSISSTRICT`, `OWNERSHIPINVARIANTINDUCTIVE`, `VERSIONMONOTONICITYLEMMA`, and `ACPLockOrderIsDeadlockFree`. Dafny 4 closes 19 verification obligations from these 9 lemmas, 0 errors. The Dafny

TABLE 3

Verification coverage of the S-BUS ACP. Each layer covers what the previous cannot. IronFleet [25] is the gold standard (full refinement, 10+ person-years); S-BUS matches standard industry practice short of IronFleet.

Tool	What is proved	Scope
TLC	ReadSetSoundness + companions	$N=3$ exhaustive (20.8M states, depth 28); $N=4$ reduced (2.8M, depth 24)
TLAPS	ReadSetSoundness + ORICCommit-Safety	Arb. N , 687 obl., 1 ax.
TLC dist.	ORI safety distributed	3 nodes, 2 agents, 247K states
Dafny	9 lemmas / 19 obligations	Concrete types, 0 errors
Empirical	0/427,308 active conflicts	$N \leq 64$, 3 CC backends

types are structurally equivalent to the Rust implementation’s; this is parallel specification, not refinement.

Distributed correctness (TLC, abstract): Module `SBus_Distributed.tla` model-checks an abstract 3-node deployment with five state variables (`registry`, `per-node delivery_log`, `leader`, `bounded term`, `per-agent last_commit_fresh`) and four transitions (`ELECTLEADER`, `AGENTGET`, `AGENTCOMMIT`, `AGENTRECOVER`). At $\text{Agents}=\{a_1, a_2\}$, $\text{Shards}=\{s_1, s_2\}$, $\text{Nodes}=\{n_1, n_2, n_3\}$, $\text{MaxVersion}=3$, with symmetry reduction over agent and node permutations: TLC explores 247,249 distinct states to depth 28 with 0 violations of `TYPEINVARIANT`, `VERSIONMONOTONICITY`, and the central `ORISAFETY` invariant. A separate temporal property `FAILOVERGAPEXISTS` confirms that the model deliberately exposes the ~ 5 ms concurrent-failover window of Limitation 11; the trace `ELECTLEADER` $\rightarrow$ `AGENTCOMMIT` with reset DL reaches the unvalidated-commit state. `SBus_Distributed` abstracts Raft itself: we rely on the standard Raft safety result [23] for log replication and assume the refinement-mapping construction composes soundly with our ORI-layer abstraction (full mechanisation deferred, Limitation 18). Exp. DR-9 (§7.21) closes the loop empirically: 30/30 ORI invariants survive injected leader failover on the deployed Rust+Raft implementation. State-variable enumeration, action transition relation, and the full reproduction artefacts are in the `sbus-formals` repository (`models/SBus_Distributed.tla`).

3.5 Observability Gap: p_{hidden} and Practical Scope

The $p_{\text{hidden}} = 0.739$ measurement (Exp. PH-2, §VII-F) is the empirical result underpinning C6. ORI’s structural guarantees apply to the 26.1% of reads in R_{obs} ; the remaining 73.9% (R_{hidden}) are not observable at the HTTP layer. This is a workload-conditional coverage bound for SWE-bench / GPT-4o-mini; different models and domains exhibit different values (domain range 51.1%–83.6%, Table 11).

Reconciliation with $f_{\text{obs}}^{\text{total}} = 0.998$ in Exp. PROXY-PH2.

Two coverage figures appear in the paper that may seem contradictory: this section reports $R_{\text{obs}} = 26.1\%$, while §7.11 reports $f_{\text{obs}}^{\text{total}} = 0.998$. These figures share neither numerator nor denominator and answer different questions. $R_{\text{obs}} = 26.1\%$ is *single-step HTTP coverage*: of all the shards that influence an agent’s reasoning at any one step, what fraction did the agent HTTP-fetch in that same step? $f_{\text{obs}}^{\text{total}} = 0.998$ is *session-cumulative DeliveryLog coverage*: of the shards an agent self-reports having used at a given step, what fraction is present in its session `DeliveryLog` at commit time — including shards GET-fetched in earlier steps and retained in DL under TTL? Concretely, if an agent fetches `models_state` at

step 5 and re-references it at step 12 from conversation memory without re-fetching, this read counts as *outside* R_{obs} at step 12 (no HTTP GET that step) but *inside* the `DeliveryLog` at step 12’s commit (the step-5 GET populated DL). The two metrics are intentional and complementary: R_{obs} is the formal-evidence regime where HTTP-recorded version is the read-time version (TLAPS-proven safety applies); $f_{\text{obs}}^{\text{total}}$ is the operational regime where session-scoped DL retains earlier HTTP-recorded versions and validates them unchanged at commit time. Section 7.11 decomposes the gap between the two into HTTP-this-step (0.443), DL-accumulation (0.555 within-row uplift), and proxy-marginal (0.0018 paired); the proof in §V (Formal Evidence) covers the first two without modification because both populate DL with the agent’s read-time version.

Formal R_{hidden} staleness bound. An agent that last updated its conversation context at step s and commits at step $s+k$ may commit with context staleness up to $\Delta = k \cdot t_{\text{step}}$ seconds. Under LLM inference with log-normal step times (mean μ , variance σ^2), the expected staleness is $k \cdot e^{\mu + \sigma^2/2}$. In our experiments ($\mu \approx \ln 5$, $\sigma \approx 0.4$, $k \leq 15$): worst-case context staleness $\leq 15 \times 5.4 \approx 81$ s per session.

3.6 Liveness with Bounded Starvation

Distinct-shard topology ($\text{SCR} = 0$): every correct agent commits on the first attempt. Shared-shard topology (Exp. E): the ACP enforces retry budget K (default $K=5$). $\text{Pr}(\text{commit within } K) = 1 - \text{SCR}^K$. At $\text{SCR} = 0.856$ ($N=16$): $K=19$ required for 95% success. Default $K=5$ is insufficient for shared-shard $N > 8$: at $\text{SCR} = 0.856$, $K=5$ gives $P(\text{success}) = 47\%$. Practitioners must set $K \geq \lceil \log(0.05)/\log(\text{SCR}) \rceil$; any finite K eliminates indefinite starvation. On `RETRYBUDGETEXHAUSTED`, the recommended agent behaviour is exponential backoff followed by fresh-session re-establishment.

Remark 6 (Correlated conflicts). *The bounded starvation analysis assumes i.i.d. conflict arrivals. Real LLM workloads exhibit bursty, positively correlated conflicts. Under positive correlation, the effective SCR rises above the measured mean, requiring $K \geq 2 \times$ the i.i.d.-derived bound. Operators in bursty environments should monitor empirical retry distributions (Limitation 9).*

4 DELIVERYLOG

The `DeliveryLog` records (k, v) for every `GET /shard/:key?agent_id=X`. At commit: $\hat{R} = R_{\text{explicit}} \cup \text{session-deliveries}$. Explicit entries take precedence. ARSI mode: agent declares its full read-set explicitly; `DeliveryLog` FP rate = 0. Default mode: `DeliveryLog` reconstructs the read-set; occasional false-positive aborts (stale entry after TTL) but zero false-negatives under A1.

5 ARCHITECTURE AND ALGORITHM

Implementation: The ACP core is ≈ 950 lines of safe Rust; the full distributed system (ACP + Raft coordination + sled persistence) is 1,679 lines of safe Rust, zero unsafe blocks (§IX). Tokio 1.44, Axum 0.8.4. Registry: `Mutex<HashMap>` for shards and tokens. WAL: `direct File::write_all()` (SIGKILL-safe). Lock ordering: `RwLock` $\rightarrow$ `TokenMutex` (single-edge graph, no cycle; verified by `AcpLockOrderIsDeadlockFree` lemma in Dafny).

TABLE 4
SDK changes required per deployment scenario.

Deployment	SDK changes	Guarantee
HTTP/1.1, single-node, no failover	Zero	Full ORI
HTTP/2 or multi-subsystem	1 commit field (ARSI)	Full ORI
Post-leader-election (distributed)	Reissue GETs on 410	ORI after recovery
Full R_{hidden} coverage	Proxy deployment	ORI over all reads

Algorithm 1 ACP Atomic Commit (write lock held lines 2–10)

Require: key k , expected version v_e , delta δ , agent α , optional R

```

1:  $\hat{R} \leftarrow \text{DeliveryLog.buildEffRS}(\alpha, k, R)$ 
2: Acquire RwLock write lock
3: for each  $(k', v') \in \hat{R}, k' \neq k$  do
4:   if  $\text{reg}[k'].v \neq v'$  then return CROSSSHARDSTALE
5:   end if
6: end for
7: if  $s_k.v \neq v_e$  then return VERSIONMISMATCH
8: end if
9:  $\text{token.insertIfAbsent}(k, \alpha)$  under token Mutex
10:  $s_k.c \leftarrow \delta; s_k.v += 1; \text{WAL.append}$ 
11: Release write lock (RAII); return OK( $s_k.v$ )

```

6 EXPERIMENTAL SETUP

Environment: All experiments: AWS Lightsail eu-west-2 instance (2 vCPU, 8 GB RAM). S-BUS server: single process, Tokio async runtime, default thread-pool. LLM backbone: GPT-4o-mini (OpenAI API, default temperature 0.3). Valid-run criteria: S-BUS healthy (HTTP 200 on `/stats`), $\geq 50\%$ steps completed, no API timeout >30 s; runs failing any criterion excluded ($<2\%$ of total).

Three-layer experiment structure: The experiments are organised along three distinct axes that prior reviewers conflated; conflation is the source of most reviewer confusion:

- 1) **ORI correctness (does the mechanism work?):** Exp. SR, CSV, E, SCALE, SJ-V5. These use shared-shard topology with ORI triggered on every conflict.
- 2) **Architecture comparison (how does parallel coordination compare to sequential?):** Exp. B, T3-A, Llama results. Distinct-shard topology; ORI never triggers. These measure the benefit of parallel coordination architecture, not ORI’s CC mechanism.
- 3) **Pure ORI effect (what does ORI add vs. baseline?):** Exp. ORI-ISOLATION. Holds architecture constant; varies only OCC enforcement. This is the clean causal test.

Experiments summary: Exp. B: SWE-bench (30 tasks; $N \in \{4, 8, 16\}$; 50 steps; GPT-4o-mini; 1,364 valid runs; distinct-shard topology, SCR = 0). Exp. SR: Direct Type-II/ R_{obs} stale-read injection (200 trials). Exp. CSV: Cross-shard validation (9,304 injections). Exp. E: Shared-shard contention (590 attempts). T3-A: Haiku-3 backbone (30 tasks, 813 runs). T3-B: Llama-3.1-8b-instant (14 tasks, 150 runs). Exp. PH-2: p_{hidden} (8,400 step-logs, 10 domains). Exp. SJ-V3: Semantic judge pilot (null; mechanistic flaw). Exp. SJ-V4: Semantic judge (1,000 runs; context diversity). Exp. SJ-V5: SCR dose-response structural validation (Table XII). Exp. MERGE: OCC vs. LLM-assisted merge (45 conflict pairs). Exp. SCALE: Contention at $N \in \{4, 8, 16, 32, 64\}$ (74,400 attempts). Exp. SEQUENTIAL: Direct wall-time measurement. Exp. ORI-ISOLATION: Pure ORI effect (959 trials, 10 domains). Exp. DEDICATED-SHARD: 600 trials, 10 domains (dedicated topol-

TABLE 5
Results roadmap. Five arcs covering twenty experimental subsections, plus two supporting subsections.

Arc	Question	Experiments
A1	Does ORI prevent structural conflicts?	Exp. B, Wall-Time, Exp. SR/CSV, Exp. Sequential, Exp. ORI-Isolation
A2	What is the coverage gap, and how is it closed?	Exp. PH-2, Exp. PH-3, PH-3 validation, Exp. Adversarial-Rhidden, Exp. PROXY-PH2
A3	Does ORI have safety parity with production OCC?	Exp. E/Scale, Exp. PG-Comparison (full, Rust-Native, Contention)
A4	Does the result generalise across LLM backbones?	Backbone Generalisation T3-A, T3-B (cross-backbone proxy results in A2)
A5	When is ORI semantically beneficial?	Exp. SJ-v4, Exp. Merge, Exp. SJ-v5, Exp. Dedicated-Shard, Exp. Shared-State
Supp.	Illustrative and supplementary evidence	Case Study (Django #11019), Exp. DR-9

ogy). Exp. SHARED-STATE: 180 trials, 3 domains (single-shard topology). Exp. DR-1...DR-9: 8 distributed sub-experiments. **Exp. PG-COMPARISON (FULL): PostgreSQL SERIALIZABLE + Redis WATCH/MULTI + S-BUS, $N \in \{4, 8, 16, 32, 64\}$, 1,350 runs, 200,880 commit attempts .**

7 RESULTS

This section presents experimental evidence in five narrative arcs followed by two supporting subsections. Each arc opens with a brief statement of its research question and the experiments that answer it; the arc order matches the contribution order in §I-D. Subsections within each arc retain their original experiment names so that cross-references throughout the paper resolve correctly. Table 5 maps each arc to its experiments.

Reading guide: A reader interested in ORI’s safety properties should read Arc A1 (structural conflicts) and Arc A3

(CC parity). A reader interested in the coverage scope should read Arc A2. A reader assessing backbone-generalisation should read Arc A4. A reader assessing deployment scope should read Arc A5; this arc is the empirical foundation for the topology-conditional contribution (C3) and motivates the adaptive-routing extension described in §11.1.

Arc A1: Structural Conflict Prevention

Arc A1 establishes that ORI prevents the structural race conditions defined in §I (Definition 1). Six experiments contribute: Exp. B compares S-BUS against three multi-agent frameworks under distinct-shard topology; Exp. SR and Exp. CSV verify the cross-shard staleness rejection mechanism directly; Exp. Sequential measures the wall-time speedup of parallel-with-ORI over sequential coordination; Exp. ORI-Isolation isolates ORI’s contribution-preservation property under contention; Exp. Workload-B extends the structural-prevention measurement to a non-code workload (data-pipeline architecture planning) across 8 domains, providing cross-distribution evidence at server-side instrumentation level.

7.1 Exp. B: Coordination Architecture Comparison (distinct-shard)

Scope: This experiment compares four multi-agent coordination architectures (S-BUS, LangGraph, CrewAI, AutoGen) on SWE-bench tasks under distinct-shard topology, in which each agent owns a dedicated shard (SCR = 0 by construction). ORI’s cross-shard stale-read rejection is never triggered in Exp. B; what is measured here is coordination-architecture efficiency, not the ORI consistency mechanism. Exp. B, Exp. T3-A, Exp. T3-B together answer “how does parallel shared-state coordination compare to sequential orchestration?”—a question distinct from “does ORI prevent SRCs?” The latter is answered by Exp. SR, CSV, E, SCALE, ORI-ISOLATION and the new Exp. PG-COMPARISON (FULL).

Coordination Fraction (CF) is a coordination-overhead metric: $CF = (\text{commit attempts} - \text{first-attempt successes}) / \text{total attempts}$. A zero-isolation system (last-write-wins) has $CF = 0$ by construction. The $46\times$ gap reported between S-BUS ($CF = 0.135$) and LangGraph ($CF = 6.18$) therefore measures the token cost of LangGraph’s workflow-orchestration overhead, not the relative quality of two CC designs. LangGraph does not implement a CC mechanism that ORI could be compared to; the apples-to-apples CC comparison is Exp. PG-COMPARISON (FULL) (§VII-N).

7.2 Task Success and Wall-Time Trade-off

AutoGen S@50 gap: Table 6 shows AutoGen achieving 88–90% S@50 vs. S-BUS 70–74%. AutoGen uses a hierarchical supervisor model with a dedicated orchestrator agent that sequences subtasks, providing implicit serialisation without OCC overhead. S-BUS targets workloads where agents must share mutable state concurrently and cannot tolerate supervisor latency; for fully decomposable tasks, AutoGen’s orchestration pattern is competitive.

Pairwise statistical tests (GPT-4o-mini): Table 7 reports Fisher’s exact tests (one-sided). S-BUS significantly outperforms CrewAI at all N ($p < 10^{-6}$); underperforms AutoGen at all N ($p < 0.01$); not significantly different from LangGraph at $N = 4, 8$ but below at $N = 16$ ($p = 0.012$).

TABLE 6

Exp. B: CF and S@50 under distinct-shard topology (SCR = 0, GPT-4o-mini, 30 tasks, 1,364 valid runs). ORI’s cross-shard rejection is never triggered here; CF measures coordination-architecture overhead, not CC mechanism quality. A no-isolation system has $CF = 0$ by construction. For ORI-triggering workloads see Table 13 (Exp. SCALE); for CC-vs-CC comparison see Table 14 (Exp. PG-COMPARISON).

System	N	CF (med)	S@50	n
LangGraph	4	6.18	70.1%	137
CrewAI	4	9.18	39.8%	103
AutoGen	4	28.1	88.0%	83
S-BUS	4	0.135	73.9%	138
LangGraph	8	6.79	75.9%	133
CrewAI	8	9.81	41.2%	102
AutoGen	8	30.6	88.6%	88
S-BUS	8	0.133	71.4%	133
LangGraph	16	7.88	82.6%	132
CrewAI	16	9.31	41.6%	101
AutoGen	16	34.2	90.2%	82
S-BUS	16	0.126	70.5%	132

TABLE 7

Pairwise S@50 (GPT-4o-mini). Fisher’s exact (one-sided). Bold: $p < 0.05$.

Comparison	N	S-BUS	Baseline	p
S-BUS vs. CrewAI	4	73.9%	39.8%	$< 10^{-6}$
S-BUS vs. CrewAI	8	71.4%	41.2%	$< 10^{-6}$
S-BUS vs. CrewAI	16	70.5%	41.6%	$< 10^{-6}$
S-BUS vs. LangGraph	4	73.9%	70.1%	0.244
S-BUS vs. LangGraph	8	71.4%	75.9%	0.203
S-BUS vs. LangGraph	16	70.5%	82.6%	0.012
S-BUS vs. AutoGen	4	73.9%	88.0%	0.006
S-BUS vs. AutoGen	8	71.4%	88.6%	0.001
S-BUS vs. AutoGen	16	70.5%	90.2%	$< 10^{-3}$

7.3 Exp. SR and Exp. CSV: Structural Validation

Exp. SR: 200/200 trials—every stale commit correctly rejected (HTTP 409), every fresh commit correctly accepted (HTTP 200). Zero errors. 95% CI upper bound: $\leq 3.0\%$.

Exp. CSV: 9,304 OCC-on injections, zero corruptions. Rule of Three 95% CI upper bound: $3/9304 = 0.032\%$. Baseline (OCC-off) corruption rate 56–62% confirms the protection is empirically necessary.

7.4 Exp. SEQUENTIAL: Measured Wall-Time Speedup

Table 8 reports directly measured wall-clock times for S-BUS parallel vs. sequential execution across $N \in \{4, 8, 16\}$ on 10 SWE-bench tasks (3 repeats per (task, N); GPT-4o-mini; 8 steps). Both conditions use the same S-BUS server, tasks, and backbone; only variable: concurrent vs. serial. Bootstrap 95% CI ($n = 2000$); Wilcoxon signed-rank one-sided.

All $p < 0.0001$ (Wilcoxon signed-rank).

7.5 Exp. ORI-ISOLATION: Mechanism Conformance Under Contention

Design: Prior experiments (Exp. B) used distinct-shard topology (SCR = 0), never triggering ORI’s rejection mechanism. Exp. ORI-ISOLATION holds architecture constant and varies only OCC enforcement. We explicitly frame this as a *mechanism conformance test*: we check that the deployed system’s retry

TABLE 8

Exp. SEQUENTIAL: measured wall-clock speedup. Near-linear scaling confirms the $\Theta(S \cdot t_{\text{step}})$ parallel execution model.

N	Parallel (s)	Sequential (s)	Speedup	95% CI
4	5.29	22.05	4.17×	[3.84, 4.40]
8	5.83	50.79	8.72×	[8.06, 9.12]
16	5.53	99.09	17.92×	[16.61, 18.90]

loop behaves in accordance with its specification under N -way contention on a shared shard. The test’s outcome is not a statistical finding about ORI’s value but an empirical confirmation that the implementation’s retry path terminates correctly and produces the specification’s admission behaviour (all contending commits eventually admitted under a finite retry budget; LWW admits exactly $1/N$).

Four conditions, $N = 4$ agents, shared-shard topology: (A) Parallel+ORI; (B) Parallel-ORI (LWW); (C) Sequential+ORI; (D) CreaAI sequential. Contention admission metric: commits/trial out of $N \times S = 4 \times 10 = 40$ intended contributions.

Remark 7 (Outcomes are specification-determined, not observed). *Both distributions are zero-variance by protocol specification: ORI’s retry loop admits every contending commit under a finite retry budget; LWW admits exactly one of N contending commits per step. LLM non-determinism affects content but not count. We therefore do not report statistical tests on this table—the numbers are conformance evidence that the deployed implementation realises the specified admission behaviour, not an effect-size measurement. Evidence of live LLM execution is wall-time variance (ORI-ON: 122–154 s; ORI-OFF: 37–45 s). Zero excluded trials.*

What this experiment establishes and what it does not: *Establishes:* the Rust implementation’s retry loop terminates correctly under N -way contention at $N = 4$, matching the specification in all 479 ORI-ON trials. This is a sanity check on the deployed implementation against the TLAPS specification, not a demonstration of ORI’s value. *Does not establish:* that ORI improves semantic output quality; that contention admission is a beneficial property in its own right. Exp. SHARED-STATE and Exp. SJ-V4 show that under single-shard topology the admitted contributions are semantically contradictory, making admission a liability rather than an asset. ORI’s value depends entirely on the topology (Exp. DEDICATED-SHARD); this experiment’s role is purely to verify that the mechanism behaves as specified when it fires, not to argue for its deployment.

Why $N = 4$: $N = 4$ is the conflict-maximal case for the contention-admission metric: under LWW, exactly $1/N = 25\%$ of contributions survive, making the ORI-ON vs. LWW admission gap maximally clear. Larger N would show smaller LWW admission rates (e.g., $1/8 = 12.5\%$); the $N = 4$ result is a conservative conformance check.

Why not PostgreSQL or Redis baselines in this experiment: This experiment is a conformance test of S-BUS’s own retry loop against its own specification; the relevant comparison is ORI-ON vs. ORI-OFF within the same deployment. Database-CC baselines (where the analogous conformance test would be PG’s own retry loop vs. its own specification) are the subject of Exp. PG-COMPARISON (FULL) (§7.13).

7.6 Exp. WORKLOAD-B: Cross-Workload Structural Validation on Data-Pipeline Planning

Motivation: Every preceding experiment in this paper measures S-BUS on SWE-bench-derived Python coordination tasks (Threat ET1, §8.1). To test whether ORI’s structural conflict-prevention mechanism generalises beyond code-coordination workloads, we constructed a non-code multi-agent workload (data-pipeline architecture planning) and re-ran the ORI-ON vs. ORI-OFF comparison with server-side instrumentation of cross-shard view-divergence at commit time.

Workload: Four agents collaboratively design a four-component data pipeline: ingestion, transformation, storage, and monitoring. Each agent owns one component shard and reads from all four. The shards have real cross-shard dependencies: the storage agent’s design must reference the transformation output format; the monitoring agent must instrument technologies the other components actually chose. Eight domains span diverse architectural pressures: e-commerce (transactional orders, eventual-consistency activity), healthcare (audit compliance, geo-replication), IoT (high ingest rate, time-series), financial (sub-millisecond latency, exactly-once), social (massive scale, hot-spot handling), gaming (high cardinality, real-time leaderboards), supply chain (heterogeneous sources, schema variation), and ad-tech (sub-50ms RTB, 30-day attribution). Each trial runs 4 agents through 4 coordination steps, with worker backbone gpt-4o-mini at temperature 0. Five trials per (domain, condition) cell yield $n=80$ trials.

Server-side instrumentation: The S-BUS server adds two atomic counters that fire under both conditions, independent of the `ori_enabled` flag. At each commit’s effective read-set evaluation, `view_checked_commits` increments; if any sibling-shard version in the agent’s read-set differs from the registry’s current version, `view_divergent_commits` also increments. Under ORI-ON the divergent commit is then rejected with HTTP-409 CROSSSHARDSTALE; under ORI-OFF the same divergence is detected and counted but the commit succeeds. This design lets us measure exactly how many stale-cross-shard commits ORI prevented, without LLM-as-judge or heuristic post-hoc analysis.

Result: Table 10 reports per-domain view-divergence rates. **Under ORI-ON:** $0/638 = 0.00\%$ **divergent commits across all 8 domains.** **Under ORI-OFF:** $590/639 = 92.33\%$ **divergent commits.** Aggregate $\chi^2 = 1094.98$, $p < 10^{-240}$. Per-domain pairwise tests are uniformly significant ($\chi^2 \in [55, 60]$, $p < 10^{-13}$ on every domain). All 80 trials reached 100% completion (every agent committed at least once); the stale-read rejections under ORI-ON were absorbed by the agent retry loop (mean 14.2–15.8 rejections per trial, recovered via re-GET-and-retry).

Interpretation and scope: This result establishes that ORI’s structural conflict-prevention mechanism fires correctly on a non-code workload across 8 architectural domains, with no false-negatives at this scale. Combined with the SWE-bench evidence in Exp. B and Exp. PG-Comparison, the C2 safety-parity claim is now supported across two distinct workload distributions.

We note three honest scope limitations. **First**, the per-domain divergence rates are uniform (91%–94%), reflecting that the workload’s structural contention pattern (4 agents writing concurrently to a shared 4-shard dependency graph) is determined by harness shape rather than domain semantics. The result demonstrates that the *mechanism* is workload-content-independent, not that ORI generalises to all multi-agent topologies. **Second**,

TABLE 9

Exp. ORI-ISOLATION: contention admission under four conditions, $N = 4$, shared-shard, 10 task domains, 959 paired trials, GPT-4o-mini. Both distributions are zero-variance by protocol specification (Remark 7). Wall-time variance (37–154 s) confirms live LLM execution.

Condition	Commits/trial	Rate	n	Wall (med.)
A: Parallel+ORI	40/40	100%	479	131 s
B: Parallel-ORI (LWW)	10/40	25%	480	38 s

TABLE 10

Exp. WORKLOAD-B: cross-shard view-divergence under ORI-ON vs. ORI-OFF on data-pipeline planning across 8 domains. “Div” = divergent commits (server-side counter, fires under both conditions); “Chk” = total commits with cross-shard read-set checked. Under ORI-ON every divergent commit is rejected with HTTP-409 and the agent retries; under ORI-OFF the divergence is detected but the commit succeeds. The 0% vs. 92% split is the structural-prevention claim made operational: ORI prevented 590 stale cross-shard commits that would otherwise have succeeded. χ^2 test on the aggregate 2×2 contingency table.

Domain	ORI-ON Div/Chk	ORI-OFF Div/Chk	ORI-OFF rate
ad-tech	0/80	73/80	91.3%
e-commerce	0/80	74/80	92.5%
financial	0/80	71/79	89.9%
gaming	0/80	74/80	92.5%
healthcare	0/78	75/80	93.8%
iot	0/80	74/80	92.5%
social	0/80	75/80	93.8%
supply chain	0/80	74/80	92.5%
Aggregate	0/638	590/639	92.3%
$\chi^2 = 1094.98, p < 10^{-240}$			

this experiment measures structural prevention (stale cross-shard reads were detected and blocked), not semantic outcome quality. Whether the 590 ORI-OFF stale commits would have produced semantically incoherent designs is not measured; an LLM-judge cross-validation was considered but rejected to avoid invoking Threat IT2. **Third**, generalisation across additional workload classes (document authoring, agent planning, retrieval-augmented generation pipelines) remains future work. The view-divergence measurement is a necessary condition for ORI to be useful on a workload, not a sufficient condition; semantic outcome-quality measurement would strengthen the claim further.

With these caveats, we judge Threat ET1 (workload distribution, §8.1) substantially mitigated: the central structural claim of this paper holds on a non-code workload at $\chi^2 = 1094.98$ significance.

Arc A2: Coverage Gap Measurement

Arc A2 measures the residual reads which fall outside R_{obs} and identifies the mechanisms that close this gap. Exp. PH-2 establishes the multi-agent baseline ($p_{\text{hidden}} = 0.739$); Exp. PH-3 evaluates ground-truth semantic extraction on a single-agent rotating-target workload; the PH-3 validation subsection examines inter-judge agreement and typed-state assumptions; Exp. Adversarial-Rhidden constructs a worst-case scenario; Exp. PROXY-PH2 decomposes structural coverage into this-step HTTP, DL-accumulation, and proxy-marginal components.

7.7 Exp. PH-2: p_{hidden} Measurement

We replicate Exp. PH across 10 task domains (Django queryset / admin / migration, Astropy FITS / WCS / units, SymPy solver / matrix, requests session, scikit-learn estimator) with 8,400 step-logs (21 tasks, 4 agents, 20 steps, 5 runs/task; GPT-4o-mini backbone).

TABLE 11

p_{hidden} by domain (Exp. PH-2; 8,400 step-logs). $f_{\text{obs}} = 1 - p_{\text{hidden}}$. Workload-conditional coverage bound on SWE-bench / GPT-4o-mini; not claimed as a universal constant.

Domain	p_{hidden}	[95% CI]	f_{obs}
Django queryset	0.511	[0.502, 0.521]	48.9%
requests session	0.575	[0.563, 0.586]	42.5%
SymPy matrix	0.612	[0.601, 0.622]	38.8%
scikit-learn	0.683	[0.674, 0.692]	31.7%
Astropy WCS	0.706	[0.697, 0.715]	29.4%
Django admin	0.732	[0.724, 0.740]	26.8%
Django migration	0.800	[0.794, 0.806]	20.0%
SymPy solver	0.810	[0.805, 0.816]	19.0%
Astropy units	0.819	[0.813, 0.825]	18.1%
Astropy FITS	0.836	[0.831, 0.841]	16.4%
Overall	0.739	[0.736, 0.741]	26.1%

Measurement methodology: At each agent step, we log every HTTP GET issued by the agent (R_{obs}) and every shard-key reference appearing in the LLM’s output text not preceded by a GET in the current session (R_{hidden}). Classification is automatic by exact string match against the registry key set. This yields a conservative lower bound on R_{hidden} : synonym references and implicit state reasoning are not counted. The true p_{hidden} is likely higher.

Result: Overall $p_{\text{hidden}} = 0.739$, 95% CI [0.736, 0.741] (Wilson score, 128,622 total reads, 95,022 R_{hidden}). Consistent with prior Exp. PH estimate of 0.706. Domain variation (Table 11): p_{hidden} ranges from 0.511 (Django queryset) to 0.836 (Astropy FITS).

7.8 Exp. PH-3: Ground-Truth Semantic Extraction and F3 Mitigation

Motivation: Exp. PH-2 (§7.7) estimated $p_{\text{hidden}} = 0.739$ on the 10-domain SWE-bench workload by classifying each shard-key reference in agent output as R_{obs} or R_{hidden} via exact string match. This measurement has two weaknesses: it measures *appearance* not *causal use*, and it provides no ground truth against which to evaluate downstream mitigation mechanisms. Exp. PH-3 addresses both by adding a self-reported ground truth and evaluating three extraction mechanisms against it.

Design: We modify the PH harness so the agent emits structured JSON `{"change": "<delta>", "shards_used": [...]}` at every step, providing an explicit list of which shards from its vocabulary were causally used. We evaluate three extraction mechanisms against this ground truth: (1) *HTTP DeliveryLog observation* (R_{obs}); (2) *keyword scan* over the agent’s last two assistant turns (the earlier `phiddenWrapper` mechanism); (3) *semantic extraction* via a dedicated analyst LLM that receives the agent’s delta, the shard vocabulary, and the fresh-content block, and returns a JSON list of causally-used shards (analyst does not receive the self-report). We run the third mechanism with three analysts in ablation: gpt-4o (cross-tier), gpt-4o-mini (same-family with worker), Claude Sonnet 4.6 (cross-family). All analysts at `temperature=0.0`.

Corpus: 2,100 step-logs across 10 SWE-bench-style domains; 21 tasks total, 5 runs per task, 20 coordination steps per run. Worker backbone gpt-4o-mini throughout. Mean self-reported shards used per step: 2.33 of 4 registered. Agent JSON parse-OK rate 100%.

Results: Table 12 reports per-mechanism recall and precision. **Keyword scan** achieves 0.073 recall at 0.520 precision: fewer than 8% of self-reported causal reads are recovered. This is a substantial revision of earlier drafts treating `phiddenWrapper` as an effective F3 mechanism; per-completion “promotion” rates measured mentions, not causal reads. **Semantic extraction** substantially outperforms keyword scan across all three analysts: gpt-4o 0.593 recall at 0.916 precision; gpt-4o-mini 0.772 at 0.791; Claude Sonnet 4.6 0.751 at 0.879 on the 630-step-log cross-family ablation. The Δ between cross-family Claude Sonnet and same-family gpt-4o-mini is -0.021 recall—the load-bearing methodological result. Cross-family convergence rejects the same-family-alignment hypothesis: an analyst that simply “matches what the worker would say” would regress to gpt-4o’s lower recall, which it does not. Semantic extraction captures a real causal signal, not a family-specific alignment artefact. Per-domain analysis (Extended TR §2) shows keyword scan reaches 0.00 recall on two of ten domains where shard functionality is referenced without naming the shard-key string; semantic extraction never drops below 0.58.

7.9 PH-3 validation summary

The PH-3 self-report ground truth was independently validated by two LLM judges (GPT-4o, Claude Sonnet 4.6) on 400 (step, shard) tasks with a frozen three-step rubric (see Appendix A for the full rubric). Both judges find that worker self-reports over-state shard usage by 29%–37%, with moderate inter-judge agreement ($\kappa=0.46$). PH-3 results should therefore be read as self-report-vs-extractor consistency rather than direct causal-read fidelity (see Threat IT1, §8.1). Disagreement concentrates at a specific rubric boundary (Step 2: “required state without direct definition”) traceable to narrative-content shards emitted by current benchmark harnesses;

TABLE 12

Exp. PH-3: three extraction mechanisms vs. self-reported ground truth. Keyword and gpt-4o rows from full 2,100-step-log evaluation; gpt-4o-mini and Claude Sonnet rows from 630-step-log cross-family ablation. HTTP recall = 1.00 on PH-3 only because agents GET every shard per step under PH-3’s design. Recall and precision are against agent self-reports; an independent validation (§7.9, Threat IT1) finds self-reports over-claim by 29%–37%.

Mechanism	Recall	Precision	F ₁
HTTP DeliveryLog	1.000	—	—
Keyword (<code>phiddenWrap.</code>)	0.073	0.520	0.128
Semantic: gpt-4o	0.593	0.916	0.720
Semantic: gpt-4o-mini	0.772	0.791	0.782
Semantic: Claude Sonnet 4.6	0.751	0.879	0.810
Combined (HTTP+keyword+sem.)	1.000	—	—

a regenerated typed-shard benchmark (§11) addresses this. Per-judge labelling outputs and the rubric prompt are released in the `sbus-experiments` repository (`run_llm_judges.py`, `score_annotations.py`).

7.10 Exp. ADVERSARIAL-RHIDDEN (summary)

We constructed an adversarial workload exposing the wrapper-layer mismatch problem: an in-process `phiddenWrapper` that detects hidden references and refreshes the `DeliveryLog` via fresh GET, but does not force LLM content regeneration, produces final-state corruption identical to ORI-OFF. The corollary is that hidden-reference mitigation must operate at the LLM-API layer, not in-process—empirical motivation for the proxy approach evaluated in Exp. PROXY-PH2 (and superseded the earlier `phiddenWrapper`, Limitation 10). Full construction (adversarial pool generation, attempt-injection protocol, content-corruption metrics, all 20/20 fail-rate trials under both ORI-OFF and the in-process wrapper, control runs under ORI-ON) is reproducible from `exp_adversarial_rhidden_v2.py` in the `sbus-experiments` repository.

7.11 Exp. PROXY-PH2: Structural-Coverage Decomposition

Exp. PROXY-PH2 runs the PH-2 multi-agent workload with a transparent LLM-API proxy and decomposes structural coverage into three components: this-step HTTP, *DL-accumulation* (session-scoped `DeliveryLog` retention of earlier HTTP GETs), and a proxy-marginal contribution. On GPT-4o-mini (16,800 paired step-logs) we measure $f_{\text{HTTP}} = 0.443$, $\text{DL-accumulation} = 0.555$, $\text{proxy-marginal} = 0.0018$ (95% CI [0.0013, 0.0024]); total $f_{\text{total}} = 0.998$. The headline finding is that ORI’s session-scoped `DeliveryLog` covers the bulk of R_{hidden} on this workload via session-scoped accumulation of prior HTTP GETs, without requiring any LLM-layer mechanism.

Cross-backbone paired replication on Anthropic Haiku 4.5 and Google Gemini 2.5 Flash ($n=2,400$ paired each) confirms safety parity (0/26,400 Type-I corruptions across all three backbones) and total-coverage conservation ($f_{\text{total}} \in [0.997, 0.999]$, $|\Delta| \leq 0.002$ between any pair). The internal attribution between this-step HTTP and DL-accumulation shifts up to $|\Delta| = 0.118$ across vendors, evidence that ORI’s coverage mechanisms are backbone-agnostic in *what* they preserve but vary in *how* coverage is decomposed. The keyword-scan proxy mechanism itself is safety-preserving but coverage-marginal and throughput-negative at realistic vocabulary sizes; semantic-extraction- at-proxy is the natural next mechanism (Limitation 16).

TABLE 13

Exp. SCALE: SCR, K_{95} , and Type-I safety at $N \in \{4, 8, 16, 32, 64\}$ (100 attempts/agent, 3 repeats; 74,400 total). Zero corruptions. Distinct-shard SCR = 0 at all N .

N	Topology	SCR	K_{95}	Corruptions
4	shared	0.676	8	0
8	shared	0.790	13	0
16	shared	0.869	22	0
32	shared	0.870	28	0
64	shared	0.747	12	0

Full per-trial coverage tables, per-backbone confidence intervals, proxy-throughput measurements at vocabulary $V \in \{4, 8, 16, 32\}$, and the structural-vs-semantic decomposition methodology are reproducible from the `exp_proxy_ph2*.py` scripts in the `sbus-experiments` repository.

Arc A3: Concurrency-Control Safety Parity

Arc A3 demonstrates that S-BUS achieves safety parity with production OCC implementations under contention. Exp. E and Exp. Scale measure shared-shard contention behaviour at $N \leq 64$; Exp. PG-Comparison implements the workload against PostgreSQL 17 SERIALIZABLE and Redis 7 WATCH/MULTI adapters and includes the Rust-Native and Contention sub-experiments controlling for adapter-language and contention-rate confounds respectively.

7.12 Exp. E and SCALE: Shared-Shard Contention at $N \leq 64$

Exp. E ($N \in \{4, 8, 16\}$): Zero Type-I corruptions across 590 commit attempts (95% CI UB $3/590 = 0.51\%$). SCR rises with N : 0.650 ($N = 4$), 0.788 ($N = 8$), 0.856 ($N = 16$). The version-check-disabled baseline shows 97.5% corruption ($N = 4$), confirming ORI’s protection is necessary and effective.

Exp. SCALE ($N \in \{4, 8, 16, 32, 64\}$, Table 13): We extend to $N = 32$ and $N = 64$. Across 74,400 total commit attempts (both topologies, 3 repeats), zero Type-I corruptions at every N (95% CI UB $3/74400 = 0.004\%$)—Property 1 holds at scale. The non-monotonic SCR decrease at $N = 64$ is a queue-theoretic pacing effect: at high agent count, Tokio write-lock serialisation imposes natural pacing on commit arrivals; as N grows past thread-pool saturation, inter-arrival time at the lock increases faster than the conflict window, compressing SCR. Zero Type-I corruptions are confirmed at all N .

Remark 8 (Liveness at scale). *At $N = 32$, $K_{95} = 28$: practitioners must set retry budget $K \geq 28$ for 95% liveness under worst-case shared-shard contention. Default $K=5$ is insufficient above $N = 8$ in shared-shard workloads. Distinct-shard deployments ($K_{95} = 1$) are unaffected.*

7.13 Exp. PG-COMPARISON (FULL): Three-Backend CC-Parity at Scale

Motivation: Reviewers across multiple rounds have flagged the absence of a database-CC baseline as the largest evaluation gap. The objection “a 50-line Redis WATCH adapter, or PostgreSQL SERIALIZABLE, would give you the same guarantee” is legitimate and cannot be dismissed on architectural grounds alone. A preliminary pilot at $N \leq 16$ (3,240 attempts) is superseded here by a full sweep at three-backend result at $N \in \{4, 8, 16, 32, 64\}$ with 200,880 commit attempts.

Setup: Three HTTP adapters expose an identical JSON

API:

- **S-BUS**: the native Rust server (port 7000). Registry is `Mutex<HashMap>`; OCC is in-process.
- **PG-SER** (`pg_sbus_server.py`, port 7001): FastAPI adapter backed by PostgreSQL 17 with `default_transaction_isolation=serializable` pinned at connection time via `conn.set_isolation_level(SERIALIZABLE)`. Each agent session maps to a database connection.
- **REDIS-WATCH** (`redis_sbus_server.py`, port 7002): FastAPI adapter backed by Redis 7 (local 127.0.0.1:6379) using `WATCH <shard_key>` followed by `MULTI ... EXEC`. Cross-shard read-set validation is implemented by WATCHING every key in the read set before the transaction block; if any version differs inside the transaction, EXEC returns null and we retry.

The harness (`pg_bench_full.py`) issues the same SWE-bench coordination workload to all three backends: 30 synthetic SWE-bench-style tasks (Django, Astropy, SymPy, SciPy, Matplotlib archetypes); $N \in \{4, 8, 16, 32, 64\}$; 3 repeats; 6 coordination steps per run. Total: **1,350 runs = 3 backends $\times$ 5 N values $\times$ 30 tasks $\times$ 3 repeats**. Each run is self-contained: the harness issues `POST /admin/reset` before and after each run and `POST /admin/shard` pre-populates shard keys. Primary metric is the count of Type-I corruptions: commits that the backend accepted despite a stale cross-shard read in the originating agent’s `DeliveryLog`.

Result: safety parity across three CC classes at $N \leq 64$: Table 14 reports the full result.

Key observations:

- **Safety (primary)**: 0 Type-I corruptions across all 200,880 commit attempts. 95% Rule-of-Three upper bound: $3/200,880 = 1.49 \times 10^{-5}$ per-attempt corruption rate (three-fold tighter than the prior Exp. SCALE bound).
- **Conflict rate**: 0 HTTP-409 conflicts were raised across all backends (discussed honestly in the following paragraph).
- **Latency**: S-BUS is faster than REDIS-WATCH by $1.80 \times - 1.95 \times$ (monotonically growing with N , consistent with Redis network-RTT per WATCH). PG-SER is $13 \times - 15.5 \times$ slower than S-BUS; PG-SER scaling exceeds REDIS-WATCH as expected (predicate lock acquisition, SSI conflict tracking, and disk fsync combine).
- **Scaling**: All three backends exhibit near-linear N -scaling ($\sim 11.5 - 13.4 \times$ from $N=4$ to $N=64$; ideal is $16 \times$). No backend exhibits super-linear degradation at $N = 64$, confirming that OCC overhead is comparable across CC classes for this workload.

Scope note: dedicated-shard produces no contention: The 0-conflict, 0-corruption outcome is expected for this workload: SWE-bench tasks use dedicated-shard topology (each of N agents owns a distinct shard), so cross-shard writes do not collide and HTTP-409 retry paths are not exercised. This experiment therefore demonstrates CC-class safety parity at scale, not CC-class effectiveness under contention; the latter is the subject of Exp. PG-CONTENTION (§).

Interpretation: The parity result rules out the informal objection “S-BUS reinvents what any DB CC system already gives you for free.” Running PG-SER or REDIS-WATCH as a substitute for S-BUS requires: (a) an HTTP adapter exposing the `GET /shard/:key + POST /commit/v2` contract (no LLM SDK speaks native SQL or Redis in a way that carries `DeliveryLog`

TABLE 14

Exp. PG-COMPARISON (FULL): three-backend CC comparison at $N \in \{4, 8, 16, 32, 64\}$. Each cell reports mean $\pm$ 1 s.d. of wall time across 90 runs (3 repeats $\times$ 30 tasks $\times$ 6 steps). **All three backends record zero Type-I corruptions and zero HTTP-409 conflicts across all 200,880 commit attempts.** 95% Rule-of-Three upper bound on the per-attempt corruption rate: $3/200,880 = 1.49 \times 10^{-5}$.

N	S-BUS (ms)	REDIS-WATCH (ms)	PG-SER (ms)	REDIS-WATCH/S-BUS	PG-SER/S-BUS
4	104.9 $\pm$ 16.2	188.9 $\pm$ 33.6	1396.7 $\pm$ 161.1	1.80	13.31
8	177.3 $\pm$ 21.3	324.0 $\pm$ 31.7	2425.2 $\pm$ 138.1	1.83	13.68
16	319.9 $\pm$ 32.6	600.8 $\pm$ 47.9	4717.3 $\pm$ 253.2	1.88	14.75
32	603.1 $\pm$ 70.5	1178.1 $\pm$ 92.5	9301.0 $\pm$ 473.3	1.95	15.42
64	1208.3 $\pm$ 110.5	2322.8 $\pm$ 208.9	18731.5 $\pm$ 1078.4	1.92	15.50
scaling ($\frac{N=64}{N=4}$)	11.5 $\times$	12.3 $\times$	13.4 $\times$	—	—

reads as first-class structured metadata); (b) a network round-trip per GET, quantified at $1.9\times$ and $15\times$ latency overhead for REDIS-WATCH and PG-SER respectively; (c) ad-hoc recovery of the automatic cross-shard read-set the DeliveryLog maintains in-server. The architectural value of S-BUS over transactional-DB baselines on this workload is therefore operational simplicity and the LLM-native contract, not structural-safety differentiation—a narrower and sharper claim than prior drafts advanced.

Adapter-language and contention extensions

Two follow-up sub-experiments control for confounds in the PG-Comparison setup. **Exp. PG-COMPARISON RUST-NATIVE** re-implements the PostgreSQL and Redis backends in matched-language Rust adapters and reruns the workload across 810 trials (136,080 commit attempts). Result: zero Type-I corruptions across all three backends; commit-throughput converges within statistical noise at $N \geq 16$ under matched adapters. The $1.6\times$ throughput advantage S-BUS shows at $N=4$ in the original PG-Comparison is a scoped effect of in-process coordination, not cross-the-board CC-mechanism superiority. **Exp. PG-CONTENTION** adds a shared-shard contention sweep (472,750 commit attempts with 427,308 active HTTP-409 conflicts) and observes behavioural parity (SCR agreement within 1pp at $N \geq 8$) plus empirical validation of the Remark 7 liveness bound within 1.5pp. Combined with the PG-Comparison main result, S-BUS, PG-SER, and REDIS-WATCH together accumulate 0 corruptions across 809,710 commit attempts spanning Python and Rust-native adapters and both dedicated- and shared-shard topologies (Rule-of-Three 95% upper bound 3.7×10^{-6}).

The architectural conclusion stands across both sub-experiments: the value S-BUS provides over transactional-DB baselines is operational simplicity and the LLM-native contract, not CC-mechanism differentiation. Per-backend latency distributions, adapter-language ablation tables, and the Remark 7 liveness-bound derivation are reproducible from `pg_bench_full.py`, `pg_comparison.py`, `pg_bench_contention.py`, and `exp_pg_contention.py` in the `sbus-experiments` repository.

Arc A4: Backbone Generalisation

Arc A4 establishes that ORI’s structural guarantees are backbone-agnostic. The T3-A and T3-B experiments replicate the multi-agent workload on Anthropic Haiku 4.5 and Google Gemini 2.5 Flash respectively, paired with the GPT-4o-mini baseline. Cross-backbone proxy results are reported in Arc A2 (Exp. PROXY-PH2).

TABLE 15

T3-B: Llama-3.1-8b-instant (Groq) backbone, 14 SWE-bench tasks, $N \in \{4, 8\}$, 3 runs/task. CWR = coord tokens / work tokens (lower = better). Mann-Whitney $p = 0.0002$ (S-BUS > LangGraph). Zero commit conflicts across all 75 S-BUS runs.

System	N	Success	CWR	Conflicts	n
S-BUS	4	84.6%	0.19	0	39
S-BUS	8	72.2%	0.18	0	36
LangGraph	4	33.3%	7.61	0	39
LangGraph	8	69.4%	9.34	0	36

7.14 Backbone Generalisation (T3-A, T3-B)

Haiku-3 (T3-A) (30 tasks, 813 runs): Under Haiku-3, S-BUS is statistically significantly below CrewAI at all N ($p < 0.01$, Fisher’s exact one-sided)—an honest negative finding. S-BUS is not significantly different from LangGraph at any N .

Llama-3.1-8b-instant (T3-B) (14 tasks, 150 runs, Table 15): Under Llama-3.1-8b, S-BUS significantly outperforms LangGraph (84.6% vs. 33.3% at $N = 4$; Mann-Whitney $p = 0.0002$).

The competitive picture is backbone-dependent: three-point spectrum: (i) Haiku-3-class (very weak): sequential coordination wins; (ii) Llama-8b (weak): S-BUS wins; (iii) GPT-4o-mini (strong): S-BUS wins. S-BUS is competitive or superior on all backbones except the weakest class where sequential information-passing dominates.

CWR (Coordination-Work Ratio): Under Llama-3.1-8b, S-BUS CWR = 0.19; LangGraph CWR = 8.44—a $44\times$ difference. LangGraph spends 8.4 tokens on coordination overhead for every 1 token of actual task work. S-BUS inverts this ratio. CWR is measured for Llama-3.1-8b-instant only; GPT-4o-mini and Haiku-3 CWR not yet measured. *Scope note*: CWR captures LangGraph’s workflow bookkeeping overhead, not CC mechanism cost; it is an architectural efficiency comparison, not CC-vs-CC.

Arc A5: Topology-Conditional Operating Envelope

Arc A5 characterises the topology-dependent semantic effect of ORI supporting Contribution C3. Exp. SJ-v4 evaluates the semantic-judge rubric under context-diversity manipulation; Exp. Merge contrasts ORI with LLM-assisted merge in the shared-shard regime; Exp. SJ-v5 measures the structural-SCR dose-response; Exp. Dedicated-Shard quantifies semantic neutrality under ORI in the dedicated-shard regime; Exp. Shared-State quantifies semantic harm under ORI in the single-shard regime. Together these establish the operating-envelope claim and motivate the adaptive-routing extension described in §11.1.

TABLE 16

Exp. SJ-v4: 20 tasks, 1,000 runs; 100% injection engagement; GPT-4o-mini judge.

Condition	n	Corrupted	Rate
Fresh (R_{obs})	500	203	40.6%
Stale (R_{hidden})	499	145	29.1%
Lift: -11.5 pp, 95% CI $[-17.4, -5.7]$ pp			
Fisher’s exact (two-sided) $p = 0.0002$			

7.15 Exp. SJ-v4: Semantic Judge Context Diversity Effect

Design: After step 5, the stale agent no longer calls GET /shard. Instead it reasons from a frozen snapshot of shard content at step 0, correctly simulating R_{hidden} . Tasks selected for cumulative-state sensitivity: 20 SymPy/Django/Astropy/scikit-learn tasks. Stale injection engaged: 499/499 stale runs (100%).

Design note: The stale agent commits at the current shard version (version check passes); its staleness is in the LLM’s reasoning context (frozen prompt), not in the version the agent presents to the ACP. ORI cannot catch this staleness: no HTTP GET for the stale content was recorded in the DeliveryLog because the agent never issued one after step 5. This is precisely the R_{hidden} problem being measured.

Result: Contrary to the hypothesis, fresh agents produced more corruption than stale agents. Fresh: 40.6%; Stale: 29.1%; lift = -11.5 pp; $p = 0.0002$. Per-task: 6/20 showed expected positive lift; 10/20 reversed negative lift; 4/20 no difference.

Interpretation: topology-driven redundancy in single-shard experiments: Exp. SJ-v4 uses a single shared shard for all four agents. In this setup, the Fresh condition causes all agents to read identical intermediate state and propose redundant patches targeting the same problem step. The stale agent, reasoning from the original problem statement, proposes early-stage work the fresh agents have already moved past, diversifying coverage. This is a topology effect, not a coordination failure.

Scope clarification: ORI is designed for dedicated-shard topology. Exp. SJ-v4 measures a degenerate case (all agents writing to one shard) that Box 2 explicitly routes away from S-BUS. In dedicated-shard topology, the Fresh condition produces non-redundant, complementary contributions (Django #11019 case study, §VII-O; all 1,364 Exp. B runs at SCR = 0). Exp. DEDICATED-SHARD (§VII-L) confirms semantic neutrality in the correct topology; Exp. ORI-ISOLATION confirms structural preservation is independent of topology.

7.16 Exp. MERGE: OCC vs. LLM-Assisted Merge

45 conflicting NL delta pairs (30 multi-domain tasks; GPT-4o-mini).

Conflict detection and recovery: OCC detected and rejected 100% of structural conflicts (45/45); 100% retry success.

Merge non-determinism: running the same merge prompt twice on identical conflicting deltas produced Jaccard word-overlap below 0.6 in 66.7% of cases (mean Jaccard = 0.544). Wilson 95% CI on the non-determinism rate at $n=45$: $[0.521, 0.786]$ (point 0.667). The interval is wide and the lower bound still $\geq 50\%$, which is the substantive claim. Scope caveat: specific to GPT-4o-mini at temperature 0.3 with one prompt design, on $n = 45$ conflict pairs; a larger-scale study with diverse models is warranted (Limitation 13).

TABLE 17

OCC vs. LLM-assisted merge trade-offs.

Property	LLM merge	S-BUS OCC
Conflict handling	Accept+resolve	Reject+retry
Extra latency	+1 LLM call	+1 retry
Determinism	Non-det. (66.7%)	Det. detection
Applicability	Resolvable NL	All NL (opaque)
Correctness	Probabilistic	Structural (R_{obs})

TABLE 18

Exp. SJ-v5: commit rates match analytic prediction exactly.

k stale	Commits	Total	Rate	Predicted
0	2,560	2,560	100.0%	100.0%
1	1,950	2,400	81.2%	81.2%
2	1,500	2,400	62.5%	62.5%
3	1,050	2,400	43.8%	43.8%

Latency: OCC median 2,361 ms (including retry); MERGE median 4,395 ms ($1.9\times$ slower).

7.17 Exp. SJ-v5: Structural SCR Dose-Response

Exp. SJ-v5 ran $N = 4$ agents with varying stale fractions $k \in \{0, 1, 2, 3\}$ (10 tasks; 57 trials per condition). The commit-rate dose-response is exact (Table 18), matching the analytic prediction $1 - (k/4) \cdot (15/20)$ exactly (prediction error < 0.5 pp). This confirms ORI’s cross-shard validation operates at the analytically predicted rate.

7.18 Exp. DEDICATED-SHARD: Semantic Quality in the Correct Topology

Motivation: Exp. SJ-v4 and Exp. SHARED-STATE (§VII-M) both show ORI harms semantic quality in single-shard tasks. A natural question: does ORI harm semantics in any topology? Exp. DEDICATED-SHARD answers by running the same agents in their intended topology: each agent owns a distinct shard. 4 agents (α_1 : ORM core, α_2 : query comp, α_3 : test writer, α_4 : reviewer), 10 SWE-bench domains, 30 runs per condition, $n = 600$.

Conditions: Fresh (ORI-ON): each agent reads current version of all other agents’ shards before committing (ORI enforces freshness via 409 rejection and retry). Stale (ORI-OFF): agents commit without read-set validation (may use stale cross-shard context). Commit rate difference (ORI-ON = 0.799 vs. ORI-OFF = 0.807) confirms ORI is structurally active in both conditions.

Result: Table 19 shows 100% coherent in both conditions across all 10 task domains. Zero redundant outputs in either condition. Lift = $+0.0$ pp—ORI is semantically neutral in dedicated-shard topology.

Why neutral: In dedicated-shard topology, agents do not compete for the same content. α_1 ’s ORM patch and α_2 ’s query patch are semantically independent—a cross-shard stale read can cause a structural 409 rejection (which ORI catches), but after retry both agents produce coherent, complementary output. The SJ-v4 / SHARED-STATE toxicity arose from agents producing semantically conflicting patches to the same shard; dedicated-shard eliminates that by construction.

TABLE 19

Exp. DEDICATED-SHARD: 4 agents, 10 task domains, $n = 600$ (300 per condition), GPT-4o-mini. Both conditions achieve 100% coherence—ORI has zero semantic cost in the correct topology.

Condition	Commit rate	Coherent	Redundant	n
Fresh (ORI-ON)	0.799	300/300 (100%)	0/300 (0%)	300
Stale (ORI-OFF)	0.807	300/300 (100%)	0/300 (0%)	300

Topology summary: Three experiments establish a consistent picture: (1) *single-shard* (SJ-v4 + SHARED-STATE): ORI is semantically harmful (forced retries accumulate contradictory patches on one shard); (2) *dedicated-shard* (this experiment): ORI is semantically neutral (100% coherence in both conditions); (3) *structural* (ORI-ISOLATION): ORI preserves $4\times$ more contributions than LWW regardless of topology. ORI’s value is *completeness of output* (all agents’ work survives) in dedicated-shard topology, not improved quality. Box 2 correctly routes single-shard tasks to sequential coordination.

7.19 Exp. SHARED-STATE: Multi-Domain Single-Shard Evaluation

Motivation: Exp. SHARED-STATE evaluates ORI on tasks with genuine single-shard contention: $N = 4$ agents compete to write to one shared shard, the use case where Box 2 already recommends sequential coordination. This experiment provides empirical evidence for that recommendation. Three domains: finance (portfolio rebalancing), healthcare (patient record update), software architecture (API schema design). $n = 180$ trials (90 per condition, 30 per domain, GPT-4o-mini).

Structural result: ORI-OFF = 0.250 (exactly $1/N$); ORI-ON = 0.534 (retry ensures most agents eventually commit).

Semantic result: Table 20 shows the re-evaluated results (corrected harness: non-empty shard seeds, real judge input). Stark: ORI-ON produces 0% consistent / 100% contradicted across all three domains; ORI-OFF produces 4.4% / 85.6%. ORI-ON is semantically worse.

Why ORI-ON is worse on single-shard tasks: ORI-ON forces all $N = 4$ agents to commit their patches via retry (commit rate 0.534 vs. 0.250). The final shard contains four sequential patches from four agents, each contradicting the previous (“equity = 65%” then “equity = 60%”...). LWW admits only one agent’s patch per step, producing less contradictory content. This is the structural reason Box 2 recommends sequential coordination for single-shard tasks: ORI’s contribution-preservation property becomes a liability when all contributions conflict semantically.

PhiddenWrapper result: We ran Exp. SHARED-STATE with the PhiddenWrapper: a Python-level intercept on every `chat.completions.create()` response that scans text for domain keywords associated with registered shard names, then issues GET to register the reference, promoting $R_{\text{hidden}} \rightarrow R_{\text{obs}}$ without agent code changes.

Result (Table 21): 500 R_{hidden} reads promoted across 480 completions (104.2% per-completion). Demonstrates feasibility of the $p_{\text{hidden}} \rightarrow 0$ path but does not validate that all promotions are causally correct—precision (fraction of promotions that are genuine causal reads) is not measured.

Cost note. The 104.2% rate does not imply doubled LLM inference cost. The wrapper performs a keyword scan ($O(\text{keywords} \times \text{tokens})$, microseconds) and at most one additional HTTP GET per detected shard reference. No extra LLM call.

Spurious promotion limitation. Keyword scanning can register incidental mentions as causal reads. An agent writing “I decided not to use the portfolio approach” would trigger a `portfolio_state` entry, potentially causing a spurious 409 rejection. ARSI mode eliminates false positives entirely; the LLM API proxy path can reduce them via semantic parsing.

Supporting Evidence

The following two subsections support claims made elsewhere in the paper but do not fit cleanly under one of the five arcs. The case study illustrates the ORI rejection mechanism on a real bug (referenced in §I); Exp. DR-9 validates the distributed-replication extension empirically (referenced in §III-D and §11.2).

7.20 Case Study: Django Issue #11019

Four role-specialised agents concurrently on Django bug #11019 (queryset ordering with `select_related()`): compiler specialist, ORM specialist, test engineer, senior reviewer. Each agent declared all four shards in its read-set on every commit. Over 15 steps (60 commit attempts), SCR = 0.0%: agents worked on dedicated shards without structural contention. Caveat: SCR = 0 means ORI’s stale-read rejection never triggered in this case study—it is a happy-path demonstration of the DeliveryLog tracking mechanism. Shared-shard contention (Exp. E) and the SJ-v5 dose-response directly validate ORI’s rejection mechanism. Final state: all four shards at $v = 15$, with consistent, non-contradictory content across compiler, ORM, test, and review components.

7.21 Exp. DR-9: ORI Survives Leader Failover

Motivation: The most-cited architectural gap across reviewer rounds: the DeliveryLog (session state) is node-local and not Raft-replicated. On leader failover, the new leader’s empty DeliveryLog means ORI cannot detect stale reads for ongoing sessions. §IX outlined a three-path solution; Exp. DR-9 validates P1 (lazy micro-batch Raft replication).

P1 mechanism: After each GET, the handler fires a `tokio::spawn` that writes a `CommitEntry::Delivery` to the Raft log (fire-and-forget: GET response returns immediately). The state machine applies `Delivery` entries on all nodes, populating their DeliveryLogs. On leader failover, the new leader’s DeliveryLog is already populated from replicated entries—no HTTP 410 recovery required.

Protocol: Each trial: (1) find current leader; (2) agent α_1 issues GET (DeliveryLog replicates via P1, ~ 5 ms); (3) kill leader; (4) wait for Raft election; (5) agent α_2 bumps shard version on new leader; (6) α_1 attempts commit with stale version on new leader. Expected with P1: HTTP 409 CROSSSHARDSTALE.

TABLE 20

Exp. SHARED-STATE: single-shard topology, 3 domains, $N = 4$, GPT-4o-mini ($n = 180$). ORI-ON preserves more contributions but produces more contradictions because all N patches conflict on the same shard. Confirms Box 2.

Condition	Commit rate	Consistent	Contradicted	n
S-BUS ORI-ON	0.534	0/90 (0%)	90/90 (100%)	90
S-BUS ORI-OFF	0.250	4/90 (4.4%)	77/90 (85.6%)	90

TABLE 21

PhiddenWrapper: $R_{\text{hidden}} \rightarrow R_{\text{obs}}$ promotion across 480 completions. Rate $> 100\%$ because completions reference multiple shards simultaneously. Keyword-scan recall with unmeasured precision; see Limitation 10.

Shard	Promotions	Domain
portfolio_state	200	Finance
patient_record	152	Healthcare
api_schema	148	Software arch.
Total / per-completion	500	104.2%

TABLE 22

Exp. DR-9: P1 session replication, 30 trials, 3-node Raft cluster. ORI holds across leader failover in 100% of trials.

Metric	Value
Trials completed	30/30 ✓
ORI held (HTTP 409)	30/30 (100%)
ORI missed (HTTP 200)	0/30 (0%)
Election time median	1,981 ms
Nodes killed	0: 13x, 1: 10x, 2: 7x
New leader	0: 11x, 1: 12x, 2: 7x

Result: Table 22. 30/30 trials: ORI held (100%). Zero trials missed. All three nodes served as leader and as killed node across the 30 trials. With $n=30$ the Wilson 95% confidence interval on the success rate is $[0.886, 1.000]$ (Rule-of-Three upper bound on the failure rate: $3/30 = 10\%$). The point estimate is 100% but the confidence interval is wide because n is small; we report DR-9 as a feasibility validation of the P1 mechanism, not as a statistically tight bound. A larger DR-9 sweep ($n \geq 300$) would tighten the lower bound to $\geq 99\%$ and is queued as future work.

Remaining limitation (concurrent failover window): If a leader fails concurrently with a GET (within the ~ 5 ms fire-and-forget replication window), the `DeliveryEntry` may not reach a majority before the new leader takes over. ORI does not hold for that specific session—equivalent to pre-P1 behaviour. The 30/30 DR-9 validation covers the sequential GET-then-kill pattern; the concurrent case is Limitation 11. In practice, LLM inference latency (median 131 s) makes the 5 ms window negligibly rare relative to the 250 ms heartbeat period.

8 LIMITATIONS

We group limitations by kind: *structural* (properties of the problem we cannot solve without redesigning ORI), *evidential* (properties claimed but not yet fully evaluated), and *mechanisation* (formal-verification gaps).

Structural

- 1) R_{hidden} **structural coverage: decomposed and partially closed by DL-accumulation.** ORI covers the 26.1% of reads in R_{obs}

structurally, with formal guarantees. The remaining $p_{\text{hidden}} = 0.739$ (domain range 51–84% on SWE-bench / GPT-4o-mini) is by definition not directly observable at the HTTP layer within a single step. Exp. PROXY-PH2 (§7.11, new) decomposes this residual into: (b) *DL-accumulation* from session-scoped DeliveryLog retention of earlier HTTP GETs (a previously-unnamed structural coverage mechanism that contributes 0.555 within-row uplift above $f_{\text{obs}}^{\text{http}}$ under PROXY-OFF on the multi-agent PH-2 workload, $n = 8,400$); and (c) pure-semantic references the agent never HTTP-fetched but that the LLM emits in context (proxy-capturable, paired marginal 0.0018 at $V=4$, decreasing to negative contribution at $V \in \{8, 12\}$). Component (b) is always structurally captured by ORI’s session-scoped DL because HTTP-populated entries carry the agent’s true read-time version; it is covered by the TLAPS safety theorems unchanged. Component (c) is captured by the transparent-proxy mechanism under skip-if-exists register semantics (Type-I = 0/16,800, RoT 95% upper bound 3.57×10^{-4}), but imposes a throughput cost scaling with vocabulary size (-2.1 pp commit rate at $V=4$; -47.2 pp at $V=12$). The residual semantic gap — proxy-captured references whose true read-time version is unknown — is the target of ongoing analyst-LLM-at-proxy work (Limitation 16). Absolute $f_{\text{obs}}^{\text{total}}$ values are conditional on self-report signal (Limitation 14); we report paired uplift (bias cancels in expectation) as the submission-grade proxy-marginal result.

- 2) **Topology restriction.** ORI is semantically neutral in dedicated-shard topology and harmful in single-shard topology (Exp. DEDICATED-SHARD vs. SHARED-STATE). Box 2 provides deployment guidance; there is no automatic topology detection.
- 3) **HTTP/2 breakage of A1.** DeliveryLog completeness relies on FIFO-per-TCP-connection ordering; HTTP/2 multiplexing can violate this. Mitigation: reverse-proxy pin to HTTP/1.1, or ARSI mode.
- 4) **Composition across subsystems (Remark 4).** ORI is not closed under subsystem composition without ARSI.

Evidential

- 5) **LLM-as-judge not validated against humans.** The semantic judge (GPT-4o-mini) has not been validated against human

experts with inter-annotator agreement on our rubric for multi-agent patch coherence. The LLM-as-judge methodology itself [41] is well-established but requires task-specific validation, which we have not performed for our semantic-coherence rubric. This is the single largest evidential gap in the semantic-quality results (Exp. SJ-V4, DEDICATED-SHARD, SHARED-STATE). A 100-item human-IAA study on a representative subset is the first planned follow-up. See Threat IT2 (§8.1) for the validity analysis.

- 6) **Conflict-proximate hypothesis partially validated.** Remark 5: 100% GET→COMMIT co-location in Exp. ORI-ISOLATION is evidence for the hypothesis in SWE-bench-style structured tasks only. Unstructured and long-horizon workloads may exhibit larger GET-to-commit gaps; not yet measured. See Threat IT3 (§8.1).
- 7) **p_{hidden} is workload-conditional.** The 0.739 figure is SWE-bench / GPT-4o-mini specific. No cross-domain generalisation claimed. Prior “first cross-domain measurement” framing deprecated.
- 8) **Backbone generalisation (partial).** GPT-4o-mini, Haiku-3, Llama-3.1-8b tested; Claude-Sonnet, Claude-Haiku-4.5, GPT-5 untested. Three-point spectrum established; cannot extrapolate.
- 9) **Correlated-conflict dynamics.** Bounded starvation analysis assumes i.i.d. arrivals. Bursty positively correlated conflicts require larger retry budgets; not measured.
- 10) **In-process `phiddenWrapper` is superseded by the transparent-proxy mechanism.** An earlier in-process `phiddenWrapper` keyword-scan wrapper is empirically weak on both axes previously documented: (a) *Coverage* is 0.073 recall on the PH-3 workload against ground-truth self-reports, order-of-magnitude below the effective deployment threshold; (b) *Layer mismatch* — Exp. ADVERSARIAL-RHIDDEN (§7.10) shows an in-process wrapper that refreshes the DeliveryLog via fresh GET on a keyword hit but does not force LLM content regeneration produces identical final-state corruption to ORI-OFF. These findings motivated the move to the LLM-API-layer proxy evaluated in Exp. PROXY-PH2 (§7.11), which replaces in-process wrappers as the structural-extraction mechanism. The PROXY-PH2 evaluation also shows that keyword-scan is insufficient *at the proxy layer* for PH-2 coverage closure (paired proxy marginal 0.0018, negative throughput contribution at $V \geq 8$); semantic-extraction-at-proxy is the correct next mechanism (Limitation 16). Keyword-scan should not be deployed as a stand-alone coverage mechanism at either layer; the finding “keyword scan is insufficient for R_{hidden} coverage” is now independently established in-process (PH-3) and at the proxy layer (PROXY-PH2).
- 11) **Session-state replication has a concurrent-failover window.** Exp. DR-9 validates P1 under sequential GET-then-kill; leader failure within the ~ 5 ms fire-and-forget window leaves an unreplicated DeliveryLog entry.
- 12) **Database-CC baseline (materially closed).** Exp. PG-COMPARISON (FULL) (Python adapters, 200,880 commit attempts), Exp. PG-COMPARISON RUST-NATIVE (matched-language adapters, 136,080 commit attempts), and Exp. PG-CONTENTION (shared-shard contention, 472,750 commit attempts with 427,308 active HTTP-409 conflicts; §) together close Limitation 12: S-BUS, PG-SER, and REDIS-WATCH achieve CC-class safety parity at $N \leq 32$ under both dedicated-shard and shared-shard regimes, across three adapter-language regimes, with 0 Type-I corruptions across 809,710 attempts combined. Exp. PG-CONTENTION additionally shows *behavioural* parity (SCR agreement within 1 pp at $N \geq 8$) and empirically validates the Remark 7 liveness bound within 1.5 pp. The matched-language sweep reveals that throughput converges across all three backends at $N \geq 16$; S-BUS’s $1.6\times$ advantage at $N=4$ is a scoped effect of in-process coordination, not cross-the-board CC-mechanism superiority. The remaining database-CC gap is a distributed-SQL adapter (YugabyteDB or CockroachDB), future work but not submission-blocking.
- 13) **Exp. MERGE sample size.** $n = 45$ conflict pairs limits generalisation of the 66.7% non-determinism finding.
- 14) **Exp. PH-3 ground truth is worker self-report; an independent inter-LLM-judge validation study was conducted and reported in §7.9.** The semantic-extraction recall and precision numbers in Table 12 are measured against gpt-4o-mini’s own JSON-structured reports of which shards it causally used. Agents may systematically over-report (mentioning considered-but-unused shards), under-report (omitting implicit state reasoning), or exhibit reasoning-vs-behaviour mismatch documented in LLM self-evaluation literature. We ran an independent validation study (400 (step, shard) tasks, two LLM judges — GPT-4o and Claude Sonnet 4.6 — with a frozen three-step rubric; §7.9). Inter-judge $\kappa=0.46$ (moderate); both judges independently find that self-reports over-claim shard usage by 29%–37%. The correct phrasing of PH-3 results is therefore “*the analyst matches the agent’s self-reported shard use at 0.59 recall, where self-report is a moderately reliable signal that over-claims by roughly one third,*” not “*the analyst correctly identifies 59% of causal reads.*” A human-annotator validation under the current harnesses would face the same rubric-boundary ambiguity observed between the LLM judges (§7.9); a more informative follow-up is a regenerated benchmark with typed shard contents (future work, §11). Cross-family analyst convergence (OpenAI and Anthropic within a few percentage points of recall) remains a necessary but not sufficient signal of ground-truth validity; it controls for model-family alignment, not for self-report accuracy.
- 15) **Workload-scope gap between PH-2 and PH-3 measurements (partially addressed).** The $p_{\text{hidden}} = 0.739$ figure comes from the multi-agent concurrent-write workload (Exp. PH-2, §7.7). The $p_{\text{hidden}} = 0.074$ figure comes from Exp. PH-3 (§7.8), a single-agent rotating-target workload that HTTP-GETs every shard per step by construction. Exp. PROXY-PH2 addresses the *structural* half of this gap via Exp. PROXY-PH2 (§7.11), which runs the PH-2 workload with instrumentation and decomposes structural coverage into HTTP / DL-accumulation / proxy-marginal components, finding that DL-accumulation alone covers 0.555 on top of $f_{\text{obs}}^{\text{http}} = 0.443$ for total 0.998 (i.e. ORI’s session-scoped DeliveryLog structurally covers the bulk of R_{hidden} on this workload without any LLM-layer mechanism). The *semantic* half of the gap — whether analyst-LLM semantic extraction transfers from PH-3 to PH-2 — is still open; the PROXY-PH2 evaluation uses keyword-scan, not analyst-LLM, because per-step analyst invocation at the proxy requires architectural work scoped as future work (Limitation 16).
- 16) **Analyst-LLM semantic extraction at the proxy layer is not evaluated in this paper.** Exp. PH-3 (§7.8) evaluates dedicated-analyst semantic extraction at 0.59–0.77 recall / 0.79–0.92 precision on a single-agent workload where $p_{\text{hidden}} = 0.074$. Exp. PROXY-PH2 (§7.11) evaluates a different,

simpler mechanism (keyword scan) on the multi-agent PH-2 workload where $p_{\text{hidden}} = 0.739$. Neither evaluation transfers semantic extraction to the PH-2 regime via the proxy layer: the proxy in this paper does not invoke an analyst LLM per call. Architecturally this is a natural extension (the proxy already has the request/response text; invoking an analyst LLM adds ~ 300 – 1000 ms latency and $\sim \$0.0002$ per agent step and could populate DL at the agent’s true read-time version by parsing the LLM context), but implementing and evaluating it requires reengineering the proxy into a semantic-capture gateway, which is scoped to future work (§11). Claims about analyst-LLM semantic extraction apply to Exp. PH-3 only; claims about keyword-scan-based structural capture at the proxy layer apply to Exp. PROXY-PH2, which finds keyword-scan necessary but not sufficient (safety-preserving, coverage-marginal, throughput-negative).

Mechanisation

- 17) **TLAPS proof retains one foundational typing axiom.** `SBus_TLAPS_v16.tla` mechanises `READSETSOUNDNESS` (recorded-read monotonicity) and `ORICOMMITSAFETY` (cross-shard equality at commit time) for arbitrary N_{agents} : 687 obligations proved by `tlapm`, 0 failed. One mathematical AXIOM is retained (`FUNTYPINGRECONSTRUCTION`, a primitive fact about TLA+’s typed-function-space construction not present in the standard `FunctionTheorems.tla` distribution); two parameter ASSUMES are retained on unspecified constants (standard TLA+ practice). The retained AXIOM is widely treated as obvious in TLA+ practice but is not mechanically discharged in the current artifact; full discharge requires Isabelle/TLA backend work, future work.
- 18) **Distributed (Raft) correctness not TLAPS-mechanised.** `SBus_Distributed.tla` model-checks an abstract 3-node model (247,249 distinct states, 0 violations, see §3.4 for details). Full Raft-TLAPS deferred (est. 6–12 person-months).
- 19) **No refinement proof to Rust implementation.** Standard industry practice short of IronFleet [25]; empirical coverage 275,280 zero-corruption attempts. Refinement via Verus [54] or Creusot [55] is blocked on async support for `tokio`-based code; future work.

8.1 Threats to Validity

We organise threats into the four standard categories used in systems-paper validity assessment: *internal* (causal claims within the experiment), *external* (generalisation beyond the experimental setup), *construct* (operationalisation of the underlying concept), and *statistical* (reliability of the quantitative inference). Threats are cross-referenced to the limitations in §VIII-A–C where their detailed mitigations are discussed.

Internal validity: Three internal threats. **(IT1) Self-report ground truth in Exp. PH-3.** Precision and recall are measured against agent JSON self-reports, which the inter-LLM-judge validation study (§7.9) finds over-state shard usage by 29%–37% under two independent judges ($\kappa=0.46$, moderate agreement in the Landis-Koch banding [45]). All recall numbers in Table 12 are therefore upper bounds on genuine causal-read recall. Readers should interpret PH-3 results as self-report-vs-extractor consistency, not data-dependency fidelity (Limitation 14). **(IT2) LLM-as-judge for semantic-quality outcomes.** The semantic judge (GPT-4o-mini in Exp. SJ-v4, DEDICATED-SHARD, SHARED-STATE) has

not been validated against human experts with inter-annotator agreement on our rubric. Known LLM-judge failure modes (position bias, length bias, anchoring on surface cues) documented by Wang et al. [49] may affect results. We have not formally tested for position or length bias in our specific rubric; the Step-2/Step-3 disagreement pattern we observe is consistent with task ambiguity rather than judge bias, but this is a hypothesis, not a proof. A 100-item human-IAA study is the principal planned follow-up. **(IT3) Conflict-proximate hypothesis untested on long-horizon workloads.** The 100% GET→COMMIT co-location observed in Exp. ORI-ISOLATION is evidence for the hypothesis on SWE-bench-style structured tasks only. Unstructured and long-horizon workloads (multi-hour agent runs, evolving plans) may exhibit larger GET-to-commit gaps that would weaken ORI’s effectiveness; not yet measured.

External validity: Three external threats. **(ET1) Workload distribution (substantially mitigated).** The principal evaluation suite is on SWE-bench-derived Python tasks across 10 domains. To test generalisation, Exp. WORKLOAD-B (§7.6) extends the structural-prevention measurement to a non-code workload (data-pipeline architecture planning) across 8 domains spanning diverse architectural pressures. Server-side instrumentation records 0/638 divergent commits under ORI-ON and 590/639 under ORI-OFF ($\chi^2 = 1094.98$, $p < 10^{-240}$). The structural mechanism fires correctly on both workload distributions. Generalisation to additional non-code workload classes (document authoring, agent planning, retrieval-augmented generation pipelines) and to semantic outcome quality (whether prevented stale commits would have produced incoherent outputs) remains future work; the structural necessary-condition is established. The single-shard collaborative- writing extension addressed by adaptive routing (§11.1) is preliminary work on a third workload distribution. **(ET2) Backbone generalisation.** The principal worker backbone is gpt-4o-mini, with gpt-4o and Claude Sonnet 4.6 used as analysts or judges; cross-backbone replication on Anthropic Haiku 4.5 and Google Gemini 2.5 Flash (Exp. PROXY-PH2) provides safety-parity evidence on two further vendors. Backbone-family ablations control for model-family alignment but not for the full task-distribution-by- backbone interaction. Three-vendor coverage establishes that the result is not GPT-specific but cannot extrapolate to all backbones. **(ET3) Workload-scope between PH-2 and PH-3.** The $p_{\text{hidden}} = 0.739$ figure comes from the multi-agent concurrent-write workload; $p_{\text{hidden}} = 0.074$ from the single-agent rotating-target workload. Exp. PROXY-PH2 addresses the structural half of this gap (DL-accumulation alone covers 0.555 on PH-2); the semantic half—whether analyst-LLM extraction transfers from PH-3 to PH-2—is unmeasured (Limitation 16).

Construct validity: Two construct threats. **(CT1) Benchmark-data construct.** Shards emitted by default CrewAI and LangGraph harnesses contain plan-narrative English rather than typed artefacts (code, SQL, schemas). Cross-labelling between shards occurs in 8.5% of records. The PH-3 and PH-2 workloads measure reads over narrative-content shards, not the typed-state regime where ORI’s safety properties have their cleanest interpretation. A regenerated typed-shard benchmark (§11) addresses this threat but has not yet been constructed. **(CT2) Structural conflict $\neq$ semantic conflict.** ORI prevents structural race conditions (Definition 1) but not semantic contradictions. Two agents may produce non-conflicting versions of the same shard structurally—committing in dependency order with fresh reads—while introducing semantic incompatibilities our system does not

detect. Section III-G makes this explicit: ORI’s safety property is a structural invariant over the HTTP-observable read projection, not a semantic-correctness guarantee. The topology-conditional contribution (C3) precisely characterises when this distinction matters in practice (Exp. SHARED-STATE).

Statistical validity: Three statistical threats. **(ST1) Zero-variance distributions.** Exp. ORI-ISOLATION reports 40/40 vs. 10/40 contributions preserved, both with zero variance by structural determinism; we do not report p -values on zero-variance data and instead report deterministic structural counts (Remark VII.2). For binomial-outcome experiments we report Wilson 95% confidence intervals; Rule-of-Three upper bounds are reported for zero-event outcomes. **(ST2) Unequal trial counts across contributions.** Contribution C2 (cross-CC parity) is supported by Exp. PG-COMPARISON at $n=884,110$ total commits (427,308 under active contention, the substantive concurrency-control test); contribution C3 (topology-conditional operating envelope) is supported by Exp. SHARED-STATE ($n=180$) and Exp. DEDICATED-SHARD ($n=600$), 2–3 orders of magnitude smaller. The C3 effects are large (100% vs. 0% pass rates with zero variance) so the small n does not preclude the claim, but readers should be aware of this asymmetry. **(ST3) Multiple-comparisons in cross-domain analysis.** The domain-by-domain p_{hidden} table (Table 11) reports 10 Wilson confidence intervals; we do not apply Bonferroni correction because we make no per-domain claims, only the aggregate $p_{\text{hidden}} = 0.739$ claim. Readers running per-domain inferences should adjust accordingly.

Aggregate validity assessment: The strongest residual threats are IT1 (self-report ground truth) and IT2 (LLM-as-judge), both addressable by human-IAA validation planned in §11. ET1 (workload distribution) was the third-strongest threat in earlier drafts; Exp. WORKLOAD-B substantially mitigates it by demonstrating structural-prevention on a non-code workload distribution at $\chi^2 = 1094.98$ significance. The regenerated typed-shard benchmark (CT1) addresses the remaining content-distribution concern. Threats specific to the formal evidence (the retained TLA+ axiom, lack of Rust-implementation refinement) are discussed under §VIII-C and addressed in Future Work (§11).

Raft Failover Window: Order-of-Magnitude Bound

The admitted ~ 5 ms concurrent-failover window (Limitation 11) represents the interval during which a fire-and-forget `DeliveryPayload` from the `get_shard` path may not have replicated before leader transition. We provide an order-of-magnitude bound on the violation rate under realistic assumptions:

$$\begin{aligned} \text{violation rate} &\approx \underbrace{\lambda_{\text{failover}}}_{\leq 10^{-5}/\text{sec}} \cdot \underbrace{w}_{5 \times 10^{-3}\text{s}} \cdot \underbrace{p_{\text{concurrent commit}}}_{\leq 0.1} \\ &\approx 5 \times 10^{-9}/\text{agent-sec.} \end{aligned}$$

Over 10^6 agent-seconds (~ 11.5 agent-days) this gives $\lesssim 5 \times 10^{-3}$ expected violations. Against the measured base-rate of ~ 131 s per LLM inference step, this is $\lesssim 7 \times 10^{-7}$ violations per agent-step. The bound is informational: the safety-critical case is not a high-rate phenomenon but a correctness question (the window is unprotected), which we address by scheduling TLAPS mechanisation of the Raft-replicated safety property as a precondition for claimed distributed soundness (Limitation 18). The current paper does not claim distributed safety in the TLAPS sense; it claims single-node safety TLAPS-proven modulo one

function-theory axiom plus 3-node abstract-Raft TLC coverage at 247,249 distinct states (§3.4).

9 DISTRIBUTION PATH AND PRACTITIONER GUIDANCE

9.1 Practitioner Decision Guide

Box 2: When to use S-BUS vs. alternatives

Use S-BUS when: (a) agents communicate via HTTP and cannot use database transactions directly; (b) shared NL state is non-commutative (conflicting agent proposals require conflict detection, not merge); (c) deployments have $N \leq 32$ agents and ≤ 4 nodes; (d) agents own dedicated shards (disjoint sub-problems)—this is the primary deployment pattern.

Use a transactional DB (PostgreSQL, CockroachDB) when: agents can issue SQL transactions directly and state is structured/typed. Exp. PG-COMPARISON (FULL) confirms PG SERIALIZABLE matches S-BUS structural safety at $N \leq 64$, with a $13\times$ – $15.5\times$ wall-time cost.

Use Redis WATCH when: agents need a high-throughput KV CC with the network-RTT cost you accept. Exp. PG-COMPARISON (FULL) confirms REDIS-WATCH matches S-BUS structural safety at $N \leq 64$, at a $1.8\times$ – $1.95\times$ wall-time cost.

Use CRDTs when: agent contributions are purely additive (append-only; no mutual exclusion).

Use sequential coordination (CrewAI) when: task success rate matters more than throughput, $N \leq 8$, and the LLM backbone is weak (Haiku-3-class). With stronger backbones (GPT-4o-mini-class), S-BUS’s parallel execution achieves competitive task success at significantly lower latency.

Use sequential coordination for single-shard collaborative tasks (all N agents writing to one key): Exp. SHARED-STATE shows ORI is semantically harmful in this topology.

9.2 Implementation

S-BUS ships a 3-node Raft-replicated implementation (1,679 lines of safe Rust, zero `unsafe` blocks; `openraft` 0.8.4). Raft [23] provides the cluster-wide serialisation point required by A3.

Raft log and shard replication: All mutating operations are submitted to the Raft leader. The leader appends a `CommitEntry` to the log, waits for majority acknowledgement, and applies the entry. Followers apply entries in log order, ensuring total commit order is identical on all nodes.

ORI correctness under Raft: The Raft leader runs full ACP validation before appending. The version check in the follower state machine serves as a serialisation safety net for concurrent commits. Because the Raft log is a total order and each entry is applied exactly once in order, Property 1 holds across all nodes.

Fault tolerance: The cluster tolerates one node failure (majority quorum: 2 of 3). Election timeout: [500, 1000] ms; heartbeat interval: 250 ms. Exp. DR-6 confirms 100/100 commits survive node 2 failure; Exp. DR-7 confirms election in 1,952 ms.

LLM API proxy path (R_{hidden} resolution): The production path to cover R_{hidden} is an LLM API proxy: a transparent HTTP proxy between agents and the LLM API that (a) intercepts completion responses and scans for shard-key references; (b) logs each reference as a `DeliveryLog` entry for the agent’s session; (c) promotes $R_{\text{hidden}} \rightarrow R_{\text{obs}}$ without any agent code change. This

TABLE 23
Exp. DR: eight distributed ORI sub-experiments (openraft 0.8.4, sled persistence; DR-8 uses 4-node cluster).

Exp.	Metric	Result
DR-1	Stale rejected / Fresh accepted ($n=200$ each)	200/200 / 200/200
DR-2	SCR ($N=4$, 2,000 commits) / Corruptions	0.0000 / 0
DR-3	Raft overhead vs. inference	<0.001%
DR-4	One-winner ($n=200$) / Cross-node corr.	200/200 / 0
DR-5	Throughput scaling (3 nodes)	1.26 $\times$, 42% eff.
DR-6	Post-failure commits / consistency	100/100 / 200=200
DR-7	Leader election / Post-election commits	1,952 ms / 50/50
DR-8	4-node convergence ($n=200$ each phase)	200/200 / 400=400
DR-9	PI session replication ($n=30$ trials)	30/30 ORI held

design requires no changes to agent code and is compatible with any LLM provider supporting HTTP. Implementation and evaluation of the proxy is planned as the next major release.

9.3 Experimental Validation (Exp. DR, 8 sub-experiments)

We deployed a 3-node S-BUS cluster (openraft 0.8.4, sled-backed persistent storage, shard creation Raft-replicated) and ran eight sub-experiments. All eight pass (Table 23).

10 CONCLUSION

S-BUS addresses Structural Race Conditions in concurrent multi-agent LLM state via the *DeliveryLog* mechanism: a server-side per-agent log of HTTP GET operations that automatically reconstructs each agent’s read-set at commit time, enabling optimistic concurrency control without agent SDK changes under HTTP/1.1. The consistency property the *DeliveryLog* provides is *Observable-Read Isolation* (ORI), a projection-based OCC over the HTTP-observable read fraction R_{obs} . The contribution is deliberately scoped: the combination of mechanism (*DeliveryLog*) and property (ORI) addresses a specific class of failures (write–write and cross-shard stale-read over R_{obs}), with explicit acknowledgement that the R_{hidden} fraction (73.9% of reads on the measured multi-agent workload) is not directly observable at the HTTP layer.

Four results support the scoped claim. *Formal evidence at three tiers*. Module `SBus_TLAPS_v16.tla` mechanises `READ-SET-SOUNDNESS` and `ORICOMMITSAFETY` for arbitrary N_{agents} (687 obligations discharged, one retained primitive TLA+ function-theory axiom); TLC exhaustively checks $N=3$ at 20.8M distinct states (depth 28, 0 violations) and a reduced $N=4$ at 2.8M states (depth 24, 0 violations); Dafny machine-checks 9 inductive soundness lemmas (19 verification obligations). *Three-backend safety parity under contention*. Across three independent CC implementations (S-BUS, PG-SER, REDIS-WATCH) the shared-shard contention experiment raised 427,308 active HTTP-409 conflicts with zero Type-I corruptions; SCR agrees within 1 pp at $N \geq 8$ and the Remark 8 liveness bound is validated within 1.5 pp at all N . Combined with dedicated-shard safety and throughput comparisons (456,802 additional zero-corruption commit attempts), total empirical coverage is 884,110 attempts across three adapter-language regimes with zero corruptions—of which the contention portion (427,308) is the relevant figure for safety-under-contention claims. *Structural-coverage decomposition on PH-2*. Exp. PROXY-PH2 (16,800

paired step-logs) decomposes R_{hidden} coverage into this-step HTTP ($f_{\text{obs}}^{\text{http}} = 0.443$), *DL-accumulation* (0.555 within-row uplift)—a property of ORI’s session-scoped *DeliveryLog* identified as the dominant coverage mechanism—and transparent-proxy marginal (0.0018 paired, 95% CI [0.0013, 0.0024]). The proxy is safety-preserving (Type-I = 0/16,800) but monotonically throughput-negative at realistic vocabulary sizes, motivating future semantic-extraction proxies. Cross-backbone paired replication on Anthropic Haiku 4.5 ($n=2,400$) and Google Gemini 2.5 Flash ($n=2,400$), via a multi-upstream extension to `sbus-proxy` that path-routes `/v1/messages` to Anthropic and `/v1beta/models/...` to Google, confirms safety parity (0/26,400 Type-I across all three vendors), total-coverage conservation ($f_{\text{obs}}^{\text{total}} \in [0.997, 0.999]$, $|\Delta| \leq 0.002$ between any pair), and commit-throughput collapse (−1.2 to −2.1 pp on all three). *Scope-delimited semantic extraction*. On the single-agent PH-3 workload ($p_{\text{hidden}} = 0.074$) a dedicated-analyst semantic extractor achieves 0.593 recall at 0.916 precision against self-report ground truth; cross-family analyst convergence rejects the same-model-family alignment hypothesis. Self-reports over-claim by 29%–37% under an independent inter-LLM-judge validation study ($\kappa=0.46$). Transfer to the multi-agent PH-2 regime remains open (Limitation 16).

The architectural value of S-BUS over transactional-DB baselines is operational simplicity and the LLM-native contract: the *DeliveryLog* reconstructs R_{obs} from HTTP GET traffic without agent-SDK changes under HTTP/1.1, and cross-shard validation inherits unchanged from ORI’s TLAPS-verified commit path. The principal remaining gaps—full Raft-TLAPS mechanisation, implementation refinement to the Rust source, and analyst-LLM semantic extraction at the proxy layer—are explicitly scoped as future work rather than claimed results.

11 FUTURE WORK

Four directions extend this work, ordered by their relationship to the contributions presented here.

11.1 Adaptive Routing for Single-Shard Topologies (A-BUS)

S-BUS’s topology-conditional operating envelope (C3, Exp. SHARED-STATE) establishes that ORI is the wrong primitive for single-shard collaborative writing. The structural preservation property which makes ORI valuable in dedicated-shard topologies—retaining every committed contribution—is counter-productive when those contributions are mutually contradictory. The natural extension is an *adaptive routing protocol* that selects between ORI and a merge-based protocol based on per-shard topology classification.

Concretely, this requires three additions to the architecture. First, agents declare a `additive_hint` (boolean) on each commit, signalling whether the delta is a region-scoped contribution that does not displace existing shard content. Second, a classifier observes the historical commit pattern per shard and classifies it as `DEDICATED`, `CONTENTEDCOMMUTATIVE`, or `CONTENTEDNON-COMMUTATIVE` based on contention rate and agent-declaration consistency. Third, a merge engine implements three strategies in chain—structural-by-region (replacing matching section headers), append-style concatenation (for grow-only logs), and an LLM-assisted semantic-merge fallback—producing a merged delta only when the oracle confirms commutativity from agent declarations.

A preliminary implementation (*A-BUS*)¹ is in active development; formal model, full evaluation methodology, and quantitative results are deferred to a companion paper rather than reported here.

This work would strengthen S-BUS’s deployment story: practitioners get ORI for the dedicated-shard regime where S-BUS’s formal proofs apply unchanged, and (when A-BUS is published) the adaptive extension for single-shard collaborative writing where ORI alone is harmful. The two regimes together would cover the workload space identified in §III.

11.2 Distributed Correctness via Raft-TLAPS Mechanisation

P1 session replication is empirically validated (Exp. DR-9, 30/30 ORI invariants survived leader failover) and TLC-checked in an abstract 3-node model. A residual ~ 5 ms concurrent-failover window within the fire-and-forget replication interval remains, and the Raft layer is not TLAPS-mechanised. Closing this gap requires composing the existing TLAPS proofs for ORI with a Raft TLAPS specification (e.g., the ironfleet-style refinement from [25]). This is the single largest mechanisation gap remaining and is the primary direction for an extended journal version of this paper.

11.3 PH-2 Semantic Extraction Transfer

The workload-scope gap between Exp. PH-2 ($p_{\text{hidden}} = 0.739$, multi-agent) and Exp. PH-3 ($p_{\text{hidden}} = 0.074$, single-agent rotating-target) leaves the transfer of semantic-extraction effectiveness to the high- p_{hidden} regime as an open question (Limitation 15). A re-run of the multi-agent workload with a dedicated-analyst extractor attached to each agent is well-defined; the cost is engineering plumbing and $\sim \$50$ of API spend per cell.

11.4 Regenerated Typed-Shard Benchmark

The PH-3 validation study (§7.9) identified that agent self-reports cannot be fully audited against traces because current benchmark shards, generated by default CrewAI and LangGraph harnesses, contain plan-narrative English rather than typed artefacts (code, SQL, schemas). A regenerated benchmark enforcing typed contents at generation time would eliminate the cross-labelling that drives the Step-2/Step-3 judge disagreement (§7.9, 8.5% of records) and permit the originally-intended two-human-annotator study.

REFERENCES

- [1] LangChain, “LangGraph,” GitHub, 2024.
- [2] J. Moura, “CrewAI,” GitHub, 2024.
- [3] Q. Wu et al., “AutoGen,” arXiv:2308.08155, 2023.
- [4] S. Hong et al., “MetaGPT,” ICLR, 2024.
- [5] G. Li et al., “CAMEL,” NeurIPS, 2023.
- [6] OpenAI, “Swarm / AG2,” GitHub, 2024.
- [7] Google, “Agent-to-Agent Protocol,” GitHub, 2025.
- [8] S. Yao et al., “ReAct,” ICLR, 2023.
- [9] O. Khattab et al., “DSPy,” arXiv:2310.03714, 2023.
- [10] Microsoft, “Semantic Kernel,” GitHub, 2023.
- [11] M. Cemri et al., “Why do multi-agent LLM systems fail?” 2025.
- [12] J. S. Park et al., “Generative Agents: Interactive Simulacra of Human Behavior,” UIST, 2023.
- [13] D. Terry et al., “Managing Update Conflicts in Bayou,” SOSP, 1995.
- [14] C. Packer et al., “MemGPT,” arXiv:2310.08560, 2023.
- [15] A. Adya, “Weak Consistency,” PhD thesis, MIT, 1999.
- [16] W. Lloyd et al., “Don’t Settle for Eventual,” SOSP, 2011.
- [17] C. Li et al., “Making Geo-Replicated Systems Fast as Possible, Consistent when Necessary,” OSDI, 2012.
- [18] Y. Sovran et al., “Transactional Storage for Geo-Replicated Systems,” SOSP, 2011.
- [19] D. Ports & K. Gritter, “Serializable Snapshot Isolation in PostgreSQL,” VLDB, 2012.
- [20] J. Corbett et al., “Spanner,” OSDI, 2012.
- [21] P. Bailis et al., “RAMP Transactions,” VLDB, 2014.
- [22] P. Bailis et al., “Coordination Avoidance in Database Systems,” VLDB, 2014.
- [23] D. Ongaro & J. Ousterhout, “Raft,” USENIX ATC, 2014.
- [24] L. Lamport, “Time, Clocks, and the Ordering of Events in a Distributed System,” CACM, 21(7), 1978.
- [25] C. Hawblitzel et al., “IronFleet,” SOSP, 2015.
- [26] J. R. Wilcox et al., “Verdi,” PLDI, 2015.
- [27] J. Zhou et al., “FoundationDB,” SIGMOD, 2021.
- [28] J. Betz, “TigerBeetle,” 2023.
- [29] S. Tu et al., “Speedy Transactions in Multicore In-Memory Databases,” SOSP, 2013.
- [30] C. Diaconu et al., “Hekaton: SQL Server’s Memory-Optimized OLTP Engine,” SIGMOD, 2013.
- [31] A. Dragojevic et al., “No Compromises,” SOSP, 2015.
- [32] P. A. Bernstein et al., “Orleans,” MSR TR-2014-41, 2014.
- [33] P. Bernstein & N. Goodman, “Concurrency Control,” ACM Surv., 1981.
- [34] H. T. Kung & J. T. Robinson, “OCC,” ACM TODS, 1981.
- [35] M. Herlihy & J. E. Moss, “Transactional Memory,” ISCA, 1993.
- [36] D. Dice et al., “TL2,” DISC, 2006.
- [37] A. Thomson et al., “Calvin,” SIGMOD, 2012.
- [38] D. Peng & F. Dabek, “Percolator,” OSDI, 2010.
- [39] R. Taft et al., “CockroachDB,” SIGMOD, 2020.
- [40] C. Jimenez et al., “SWE-bench,” ICLR, 2024.
- [41] L. Zheng et al., “LLM-as-a-Judge,” NeurIPS, 2023.
- [42] Loro Team, “Loro: Movable Tree CRDT,” <https://loro.dev>, 2024.
- [43] M. Kleppmann et al., “Automerger,” <https://automerger.org>, 2024.
- [44] R. Fielding and J. Reschke, “Hypertext Transfer Protocol (HTTP/1.1): Conditional Requests,” IETF RFC 7232, June 2014.
- [45] J. R. Landis and G. G. Koch, “The Measurement of Observer Agreement for Categorical Data,” *Biometrics*, vol. 33, no. 1, pp. 159–174, 1977.
- [46] Temporal Technologies, “Workflow Update API,” Temporal documentation, 2024. <https://docs.temporal.io/workflows#update>
- [47] A. Beamer et al., “FoundationDB Directory Layer,” FoundationDB documentation, 2024. <https://apple.github.io/foundationdb/developer-guide.html#directories>
- [48] C. Packer et al., “Letta: Stateful Agents Beyond Context Windows,” GitHub, 2024. <https://github.com/letta-ai/letta>
- [49] P. Wang et al., “Large Language Models are not Fair Evaluators,” arXiv:2305.17926, 2023.
- [50] A. Dey, A. Fekete, R. Nambiar, U. Röhm, “Scalable Transactions across Heterogeneous NoSQL Key-Value Data Stores,” *PVLDB*, vol. 6, no. 12, pp. 1434–1439, 2013.
- [51] D. Gao et al., “AgentScope: A Flexible yet Robust Multi-Agent Platform,” arXiv:2402.14034, 2024.
- [52] G. Wang et al., “Voyager: An Open-Ended Embodied Agent with Large Language Models,” *Transactions on Machine Learning Research*, 2024.
- [53] J. Yang et al., “SWE-agent: Agent-Computer Interfaces Enable Automated Software Engineering,” *NeurIPS*, 2024.
- [54] A. Lattuada et al., “Verus: Verifying Rust Programs using Linear Ghost Types,” *OOPSLA*, 2023.
- [55] X. Denis, J.-H. Jourdan, C. Marché, “Creusot: A Foundry for the Deductive Verification of Rust Programs,” in *Formal Methods: 25th Intl. Symp.*, 2023.
- [56] Z. Jia, E. Witchel, “Boki: Stateful Serverless Computing with Shared Logs,” *SOSP*, 2021.
- [57] K. Honda, V. T. Vasconcelos, M. Kubo, “Language Primitives and Type Discipline for Structured Communication-Based Programming,” *ESOP*, 1998.

APPENDIX

The rubric used by both LLM judges in the Exp. PH-3 validation study (§7.9) is reproduced verbatim below. The prompt was frozen before observing inter-judge agreement and was not revised.

```
You are a strict code auditor. Decide
whether the Candidate Shard provided
content that the agent demonstrably
needed to produce the Code Change.
```

1. <https://github.com/sajjadanwar0/abus>

DECISION PROCEDURE (apply in order;
stop at the first that fires):

Step 1 -- Direct entity definition.

Does the shard's content *define* (not merely mention) the specific function, class, variable, field, or constant that is being modified or read by the Change? If yes → `<label>Yes</label>`.

Step 2 -- Required state or schema.

Does the Change transform a value, structure, or invariant whose concrete shape is only recoverable from this shard's content (e.g. a schema, a prior version, a signature, a type)? If yes → `<label>Yes</label>`.

Step 3 -- Default. Topical overlap ("both relate to the database"), shared vocabulary, or mere availability in the context is NOT sufficient. → `<label>No</label>`.

HARD RULES (override everything above):

R1. If the shard content shown is empty, truncated, or does not actually contain the entity referenced in Step 1/2 → `<label>No</label>`.

R2. If you cannot quote specific words or tokens from the shard that the Change depends on → `<label>No</label>`.

R3. Mere name collisions (two shards both contain the word "user") do not count as evidence. The evidence must be semantic.

Both judges were called with `temperature=0`, `max_tokens=512`. Output-format failures were 0/0 across the 400 tasks (§7.9).